

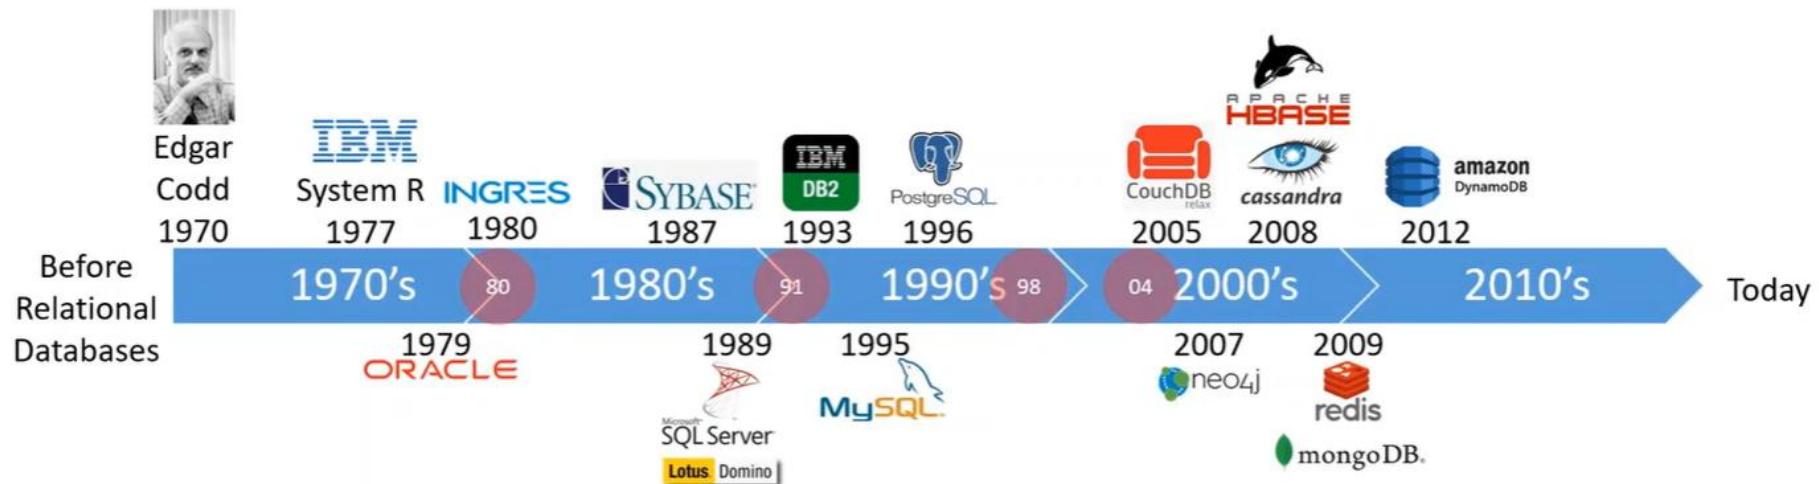
Lesson 3

DATABASES



Today's requirements on databases

- Big data (large datasets)
- Agility
- Unstructured/ semi structured data



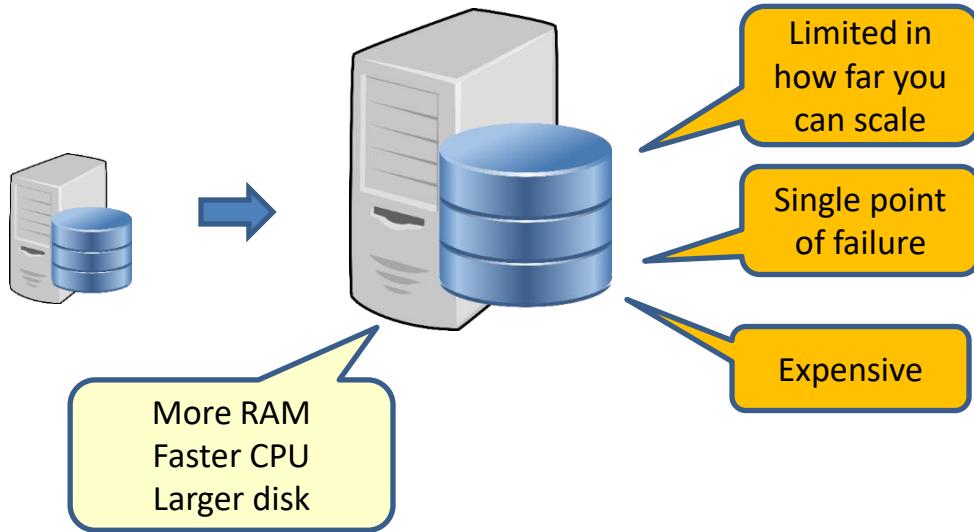
Large datasets

- Too much data
 - The data does not fit anymore on one node

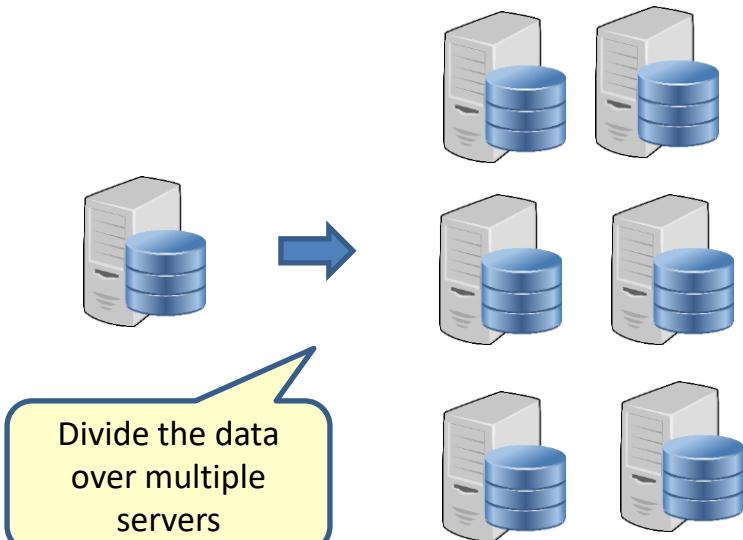


Database Scaling

- Vertical scaling



- Horizontal scaling



Five different types of databases

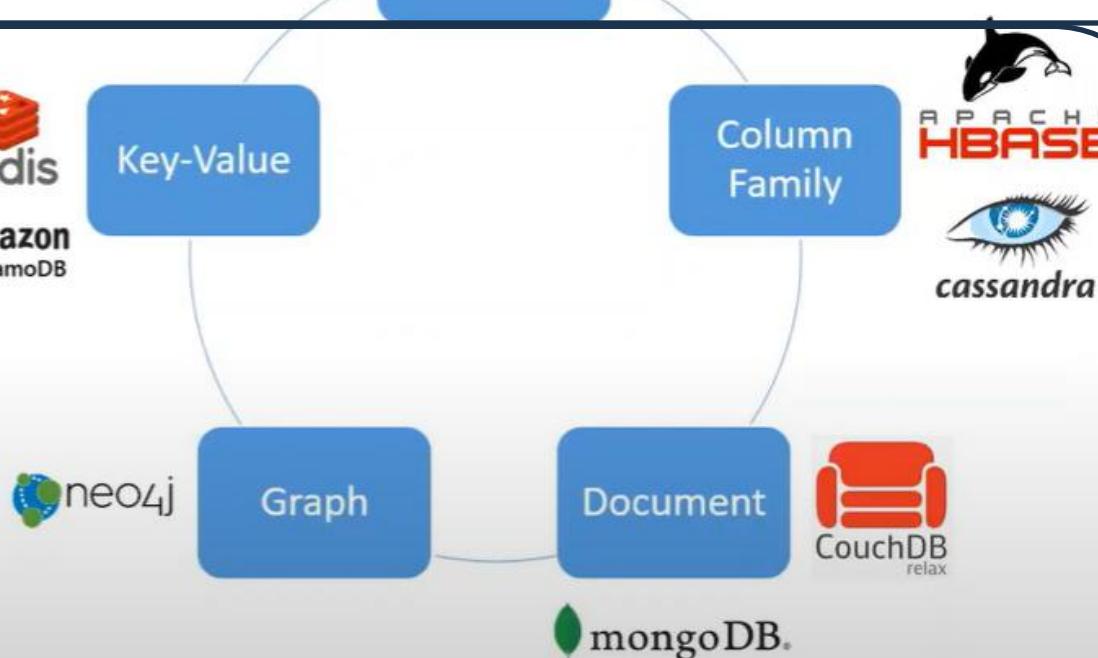
Relational



Non-relational

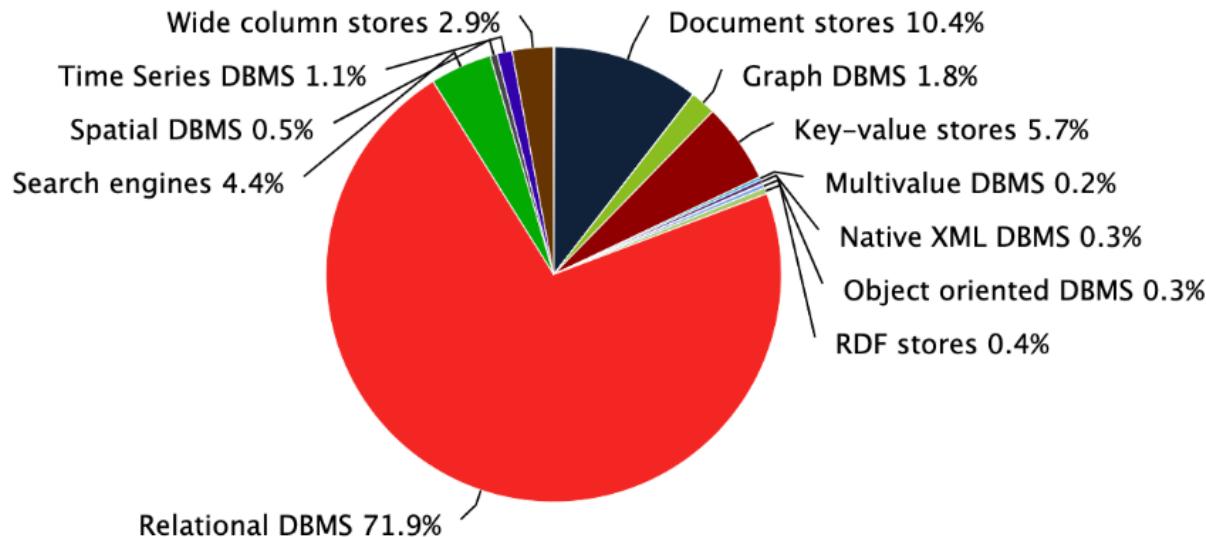


NoSQL
(Not only SQL)



Relational vs non-relational popularity

Ranking scores per category in percent, January 2023



© 2023, DB-Engines.com



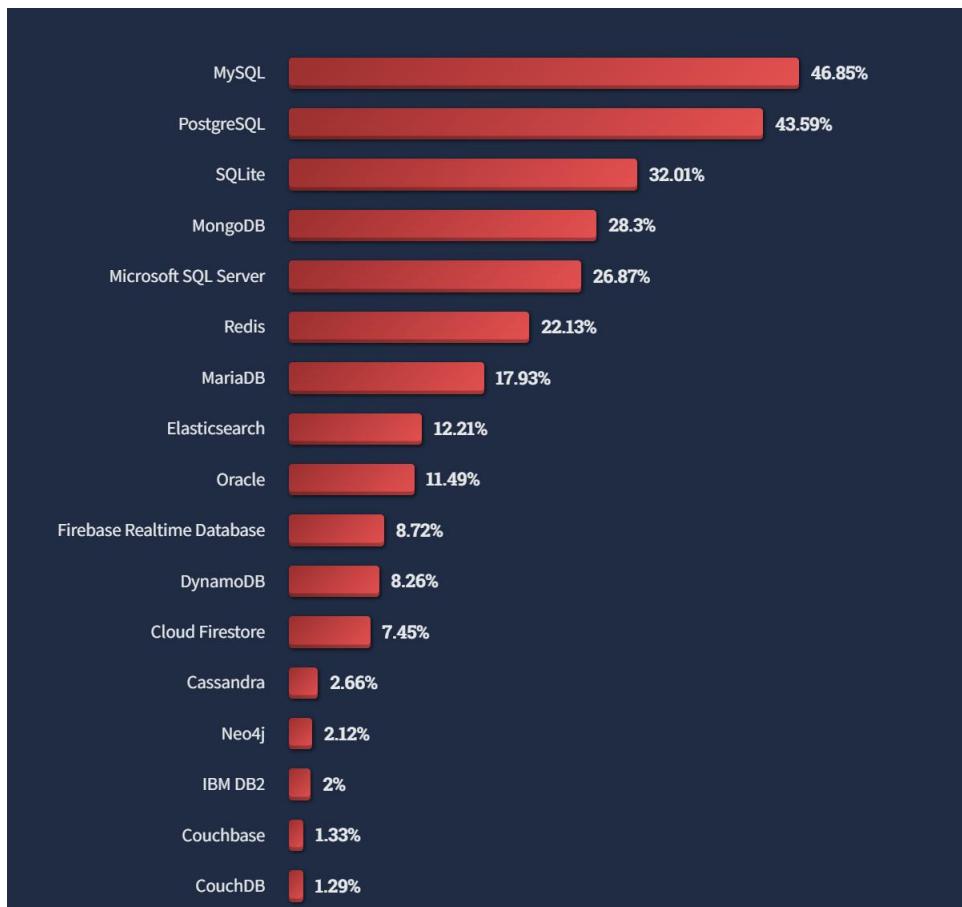
Most popular databases

<https://db-engines.com/en/ranking>

Rank May 2024	Rank Apr. 2024	Rank May 2023	DBMS	Database Model	Score		
					May 2024	Apr. 2024	May 2023
1.	1.	1.	Oracle	Relational, Multi-model	1236.29	+2.02	+3.66
2.	2.	2.	MySQL	Relational, Multi-model	1083.74	-3.99	-88.72
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model	824.29	-5.50	-95.80
4.	4.	4.	PostgreSQL	Relational, Multi-model	645.54	+0.49	+27.64
5.	5.	5.	MongoDB	Document, Multi-model	421.65	-2.31	-14.96
6.	6.	6.	Redis	Key-value, Multi-model	157.80	+1.36	-10.33
7.	7.	↑ 8.	Elasticsearch	Search engine, Multi-model	135.35	+0.57	-6.28
8.	8.	↓ 7.	IBM Db2	Relational, Multi-model	128.46	+0.97	-14.56
9.	9.	↑ 11.	Snowflake	Relational	121.33	-1.87	+9.61
10.	10.	↓ 9.	SQLite	Relational	114.32	-1.69	-19.54
11.	11.	↓ 10.	Microsoft Access	Relational	104.92	-0.49	-26.26
12.	12.	12.	Cassandra	Wide column, Multi-model	101.89	-1.97	-9.25
13.	13.	13.	MariaDB	Relational, Multi-model	93.21	-0.60	-3.66
14.	14.	14.	Splunk	Search engine	86.45	-2.26	-0.19
15.	↑ 17.	↑ 18.	Databricks	Multi-model	78.61	+2.28	+14.66
16.	↓ 15.	16.	Microsoft Azure SQL Database	Relational, Multi-model	77.99	-0.41	-1.21
17.	↓ 16.	↓ 15.	Amazon DynamoDB	Multi-model	74.07	-3.50	-7.04
18.	18.	↓ 17.	Hive	Relational	61.17	-1.41	-12.44
19.	19.	↑ 20.	Google BigQuery	Relational	60.38	-1.52	+5.51
20.	20.	↑ 21.	FileMaker	Relational	48.20	-1.53	-3.80
21.	21.	↓ 19.	Teradata	Relational, Multi-model	45.33	-2.52	-17.39
22.	22.	↑ 23.	SAP HANA	Relational, Multi-model	44.69	-1.14	-5.67
23.	23.	↓ 22.	Neo4j	Graph	44.46	-0.01	-6.65

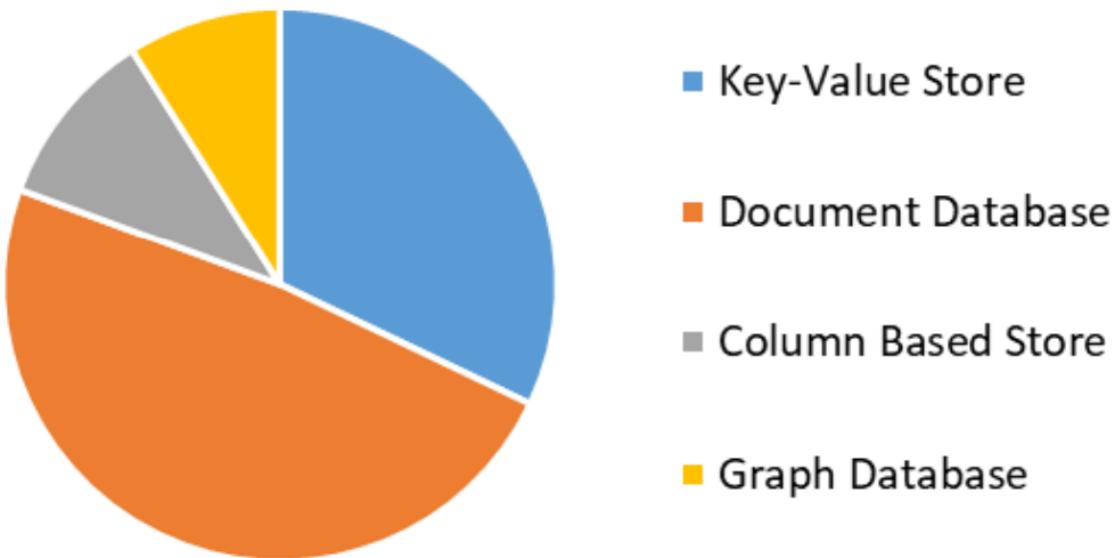
Most popular databases

<https://survey.stackoverflow.co/2022#technology>



Most popular NoSQL databases

NoSQL Database Market, by Type 2022 (%)

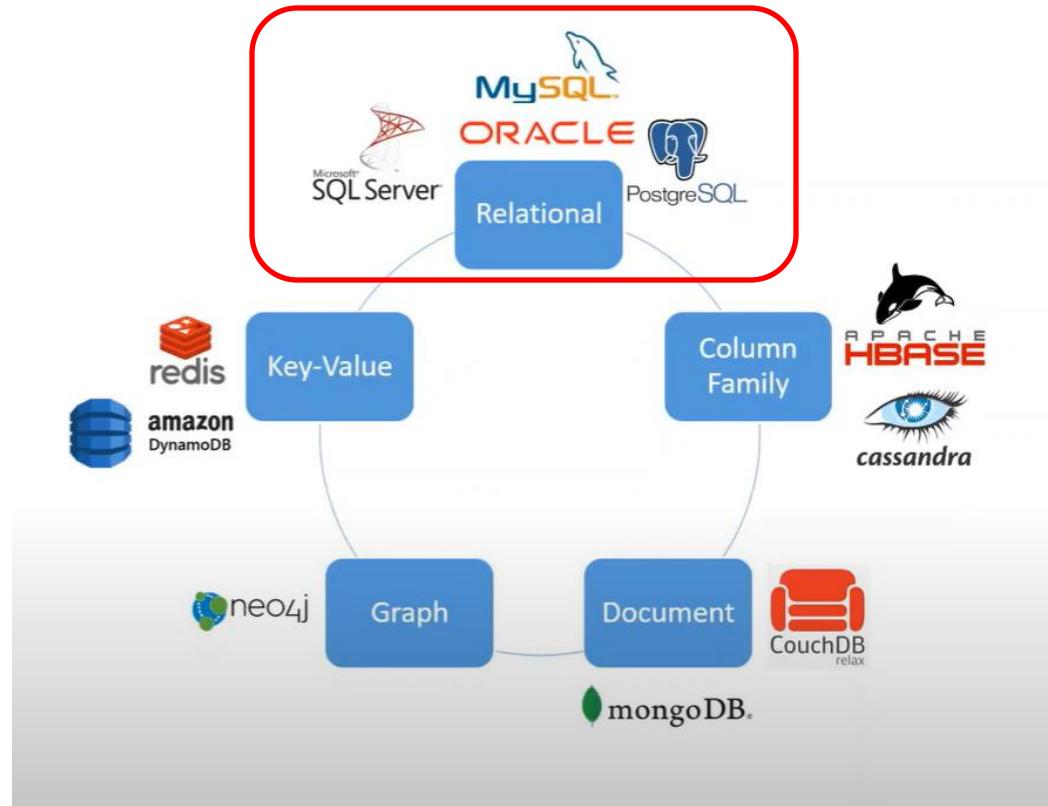


Cloud database

Cloud Database Cheat Sheet

 blog.bytebytego.com

DB Type		aws	Azure	Google Cloud	Open Source / 3rd Party	
Structured	Relational	 RDS	 SQL Database	 Cloud SQL	 Oracle	 PostgreSQL
	Columnar	 Redshift	 Synapse Analytics	 BigQuery	 MySQL	 SQL Server
	Key Value	 DynamoDB	 Cosmos DB	 BigTable	 Snowflake	 Click House
	In-Memory	 ElastiCache	 Azure Cache for Redis	 Memory Store	 Redis	 Scylla
	Wide Column	 Keyspaces	 Cosmos DB	 BigTable	 Redis	 Memcached
	Time Series	 Timestream	 Time Series Insights	 BigTable	 Cassandra	 Scylla
	Immutable Ledger	 Quantum Ledger DB	 Confidential Ledger	 CloudSpanner	 Influx	 OpenTSDB
	Geospatial	 Keyspaces	 Cosmos DB	 BigTable	 Hyper Ledger Fabric	
	Graph	 Neptune	 Cosmos DB	 CloudSpanner	 PostGIS	 geomesa
	Document	 Document DB	 Cosmos DB	 FireStore	 OrientDB	 Dgraph
Semi Structured	Text Search	 OpenSearch	 Cognitive Search	 CloudSearch	 MongoDB	 Couchbase
	Blob	 S3	 Blob Storage	 Cloud Storage	 Elasticsearch	 Elassandra
UnStructured					 Ceph	 OpenIO



RELATIONAL DATABASE

Relational databases

Data is normalized:
No duplication of
data

Key

Relationships are
expressed with
key's, foreign keys
and link tables

Relational - Tables

Customer ID	First Name	Last Name	City
0	John	Doe	New York
1	Mark	Smith	San Francisco
2	Jay	Black	Newark
3	Meagan	White	London
4	Edward	Daniels	Boston

Account Number	Branch ID	Account Type	Customer ID
10	100	Checking	0
11	101	Savings	0
12	101	IRA	0
13	200	Checking	1
14	200	Savings	1
15	201	IRA	2

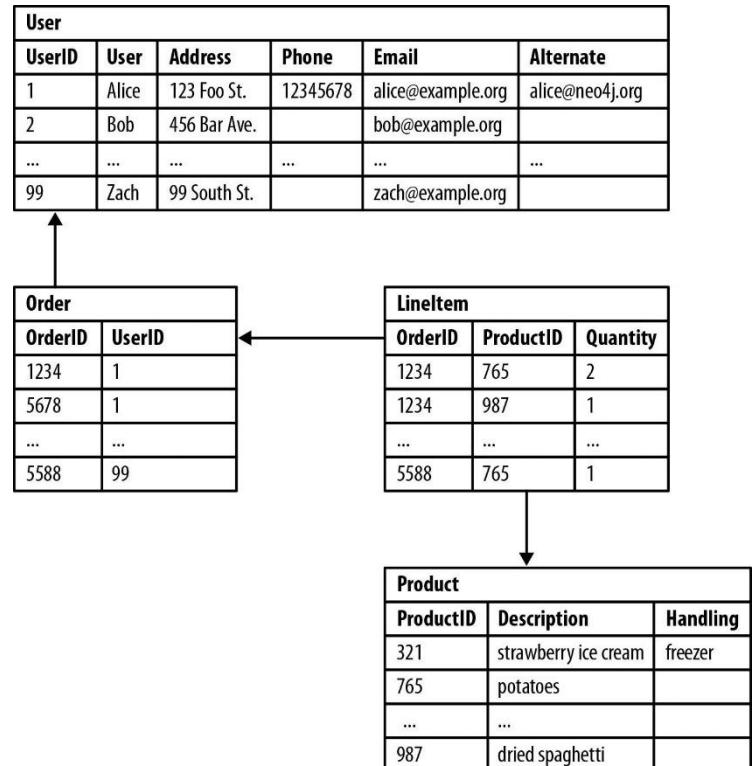
SQL as language to
manage data in
the database

Foreign key

To find related
data you need to
JOIN tables
together

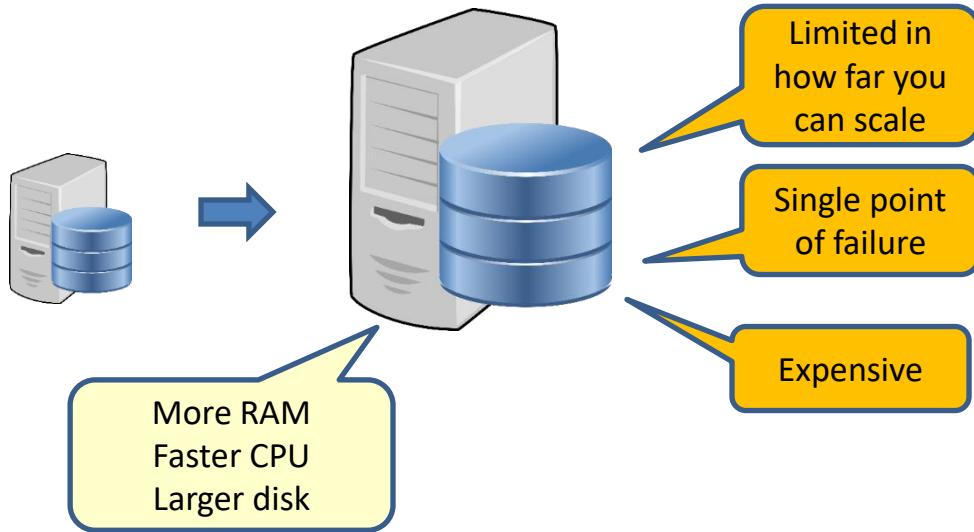
Relational database characteristics

- Structured data
- Normalized
 - No duplication
- ACID transactions
- Fixed schema
- SQL as query language
- Scale vertically
 - Difficult to scale horizontally

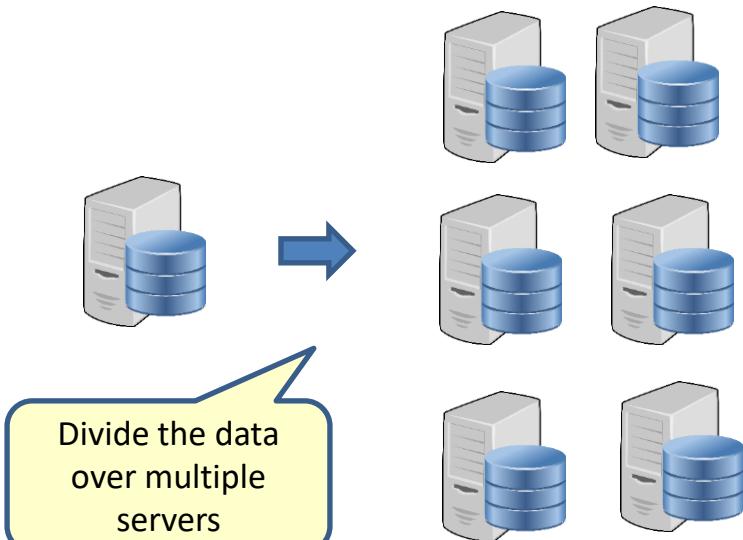


Database Scaling

- Vertical scaling



- Horizontal scaling



Horizontal scaling

- 2 techniques
 1. Sharding/partitioning: Divide the data over the database instances
 2. Replication: make a copy of the data on different database instances



Sharding



Original Table

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN
3	SELDÄ	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE

HP1

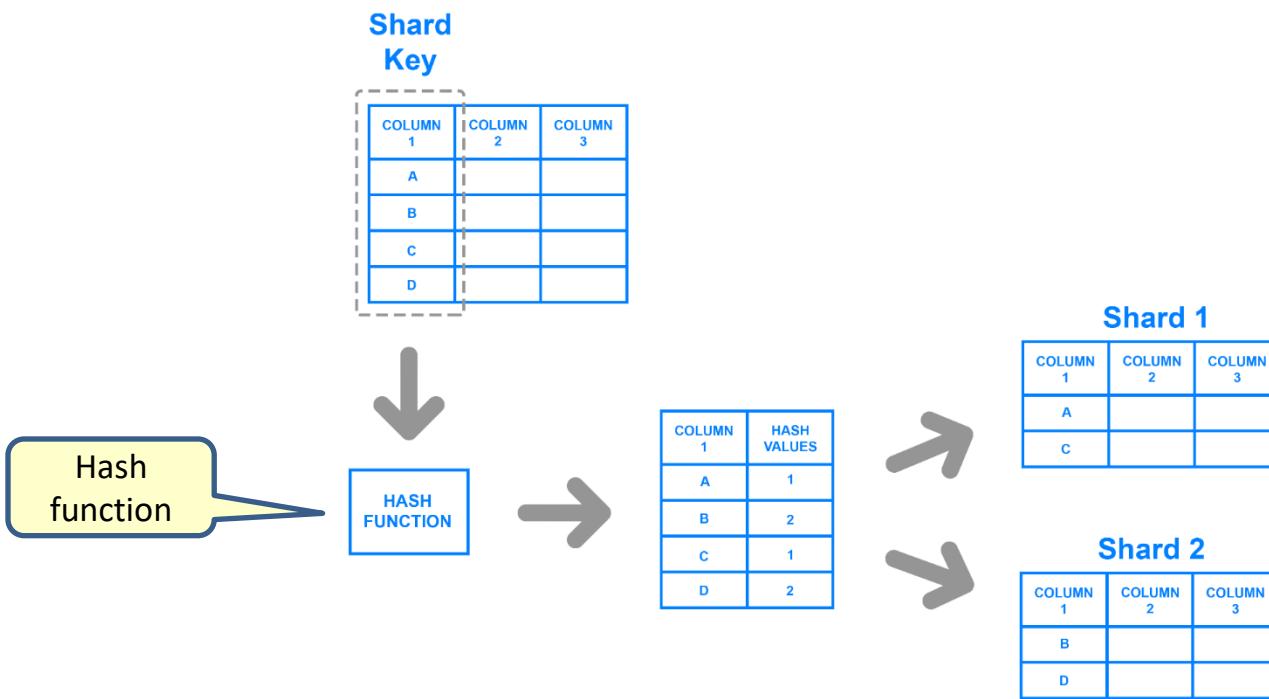
CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
1	TAEKO	OHNUKI	BLUE
2	O.V.	WRIGHT	GREEN

HP2

CUSTOMER ID	FIRST NAME	LAST NAME	FAVORITE COLOR
3	SELDÄ	BAĞCAN	PURPLE
4	JIM	PEPPER	AUBERGINE



Key based sharding



Range based sharding



PRODUCT	PRICE
WIDGET	\$118
GIZMO	\$88
TRINKET	\$37
THINGAMAJIG	\$18
DOODAD	\$60
TCHOTCHKE	\$999



(\$0-\$49.99)

PRODUCT	PRICE
TRINKET	\$37
THINGAMAJIG	\$18



(\$50-\$99.99)

PRODUCT	PRICE
GIZMO	\$88
DOODAD	\$60

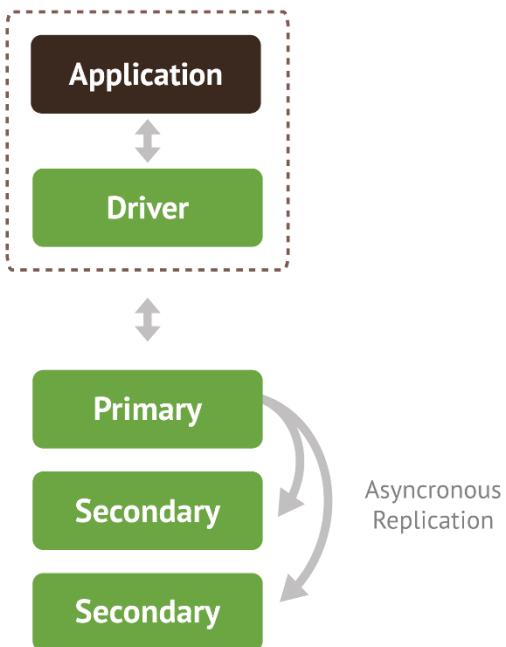


(\$100+)

PRODUCT	PRICE
WIDGET	\$118
TCHOTCHKE	\$999



Replica Sets



- Replica Set – two or more copies
- Failover
 - “Self-healing” shard
- Addresses many concerns:
 - High Availability
 - Disaster Recovery
 - Maintenance

Relational database advantages

- Data consistency and integrity
 - ACID transactions
- Data accuracy
 - No duplication
- Standard query language (SQL)
- Years of experience
- Structured database design
 - Normalization



Problems with relational databases

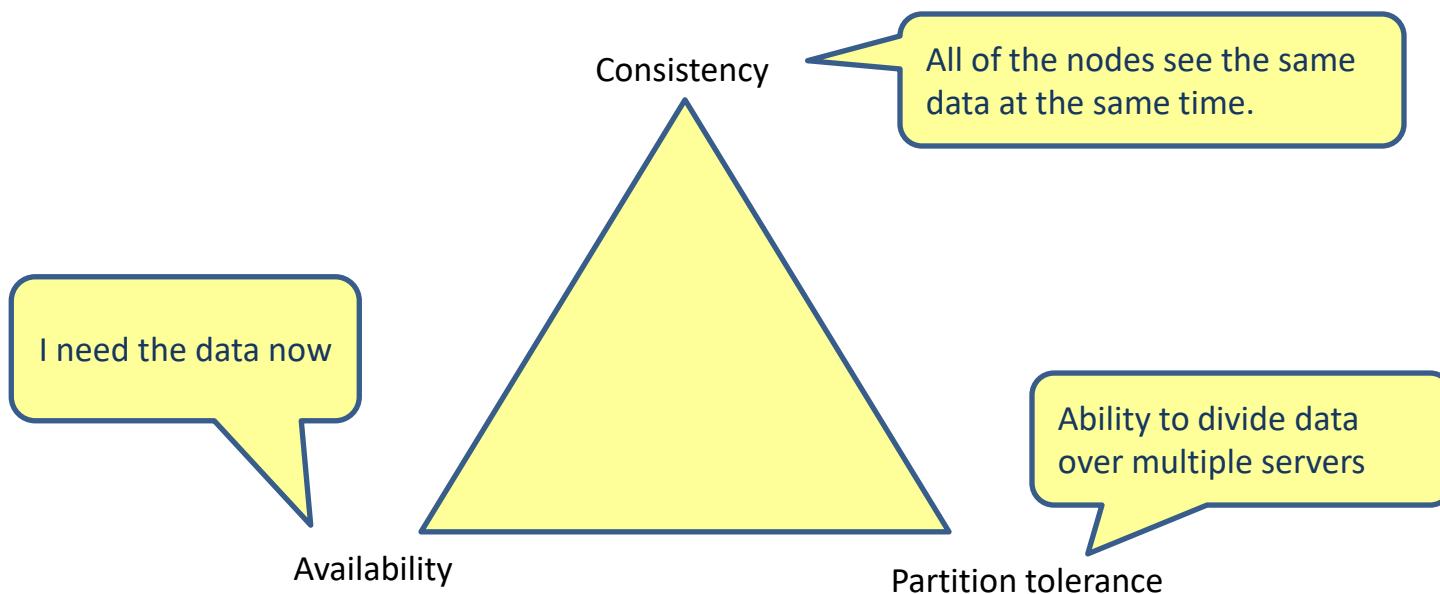
- Scaling writes are very difficult and limited
 - Vertical scaling is limited and is expensive
 - Horizontal scaling is limited and is complex
 - Queries work only within shards
 - Strict consistency and partition tolerance leads to availability problems

A relational database is hard to scale



Brewer's CAP Theorem

- A distributed system can support only two of the following characteristics



Consistency

- Strict consistency
 - The data that I read is always correct
 - You never loose data
- Eventual consistency
 - The data might not be correct
 - But will eventually become correct



Problems with relational databases

- The schema in a database is fixed
- Schema evolution
 - Adding attributes to an object => have to add columns to table
 - You need to do a migration project
 - Application downtime ...

A relational database is hard to change



Problems with relational databases

- Relational schema doesn't easily handle unstructured and semi-structured data
 - Emails
 - Tweets
 - Pictures
 - Audio
 - Movies
 - Text

Unstructured data				Semi-structured data				Structured data					
ID	Name	Age	Degree	ID	Name	Age	Degree	ID	Name	Age	Degree		
1	John	18	B.Sc.	<University>	<Student ID="1">	<Name>John</Name>	<Age>18</Age>	<Degree>B.Sc.</Degree>	</Student>	2	David	31	Ph.D.
2	David	31	Ph.D.	<Student ID="2">	<Name>David</Name>	<Age>31</Age>	<Degree>Ph.D. </Degree>	</Student>	3	Robert	51	Ph.D.	
3	Robert	51	Ph.D.					4	Rick	26	M.Sc.	
4	Rick	26	M.Sc.	</University>					5	Michael	19	B.Sc.	
5	Michael	19	B.Sc.										

A relational database does not handle unstructured and semi structured data very well



Relational database disadvantages

- Schema evolution is difficult
- Horizontal scaling is difficult
- Does not handle unstructured data very well
 - Not good in full text search
- Queries can be slow due to joins



Non relational databases

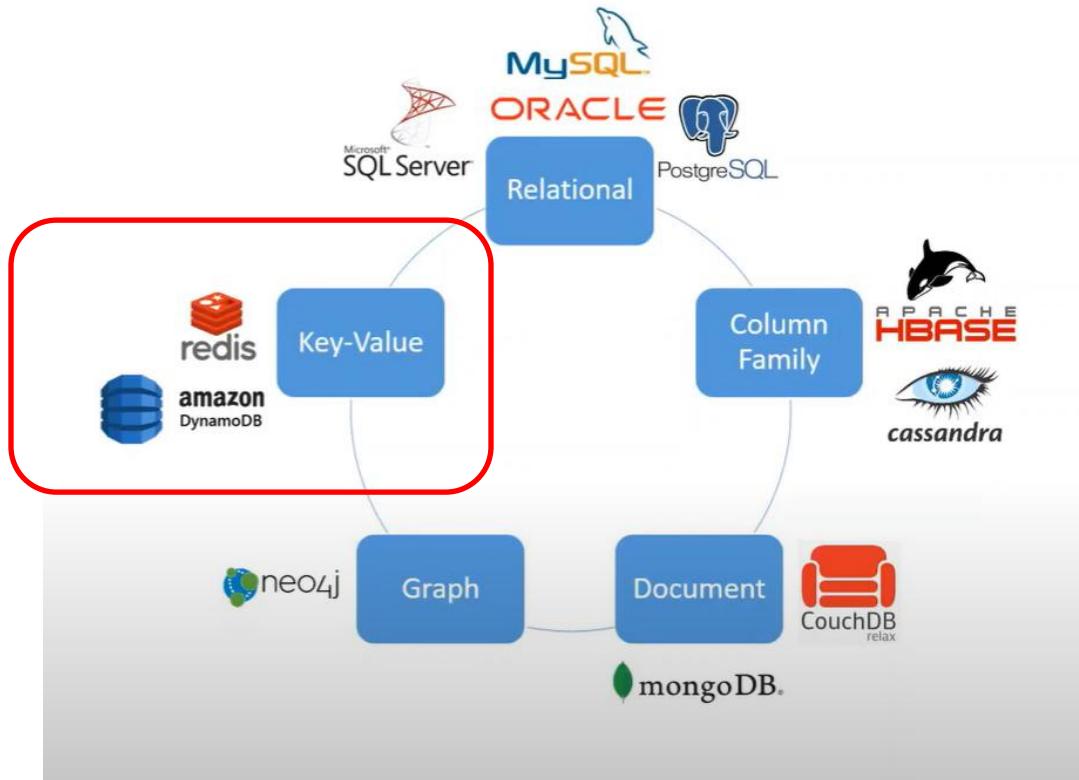
- Schema free
- Distributed
 - Horizontal scalable
 - Replicated
- Some support unstructured or semi-structured data
- BASE



BASE

- Basically available
 - All users can concurrently access the data without having to wait
- Soft state
 - Data can have transient or temporary states that change over time
- Eventual consistency
 - The data will become consistent eventually

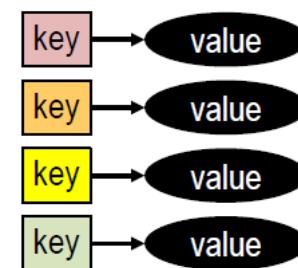




KEY VALUE STORE

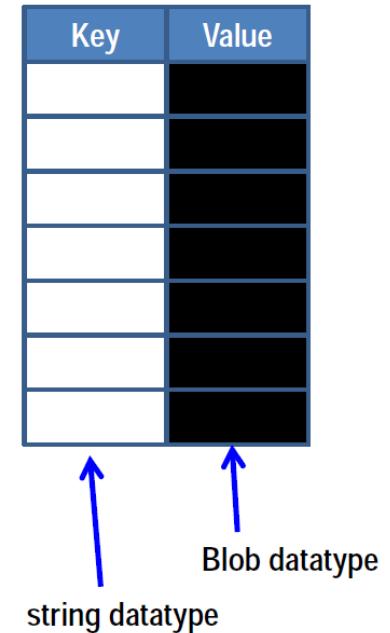
Key-value store

- One key -> one value



Key-value store

- One key -> one value
 - Simple hash table
 - Very fast
- Value is a binary object (BLOB)
 - DB does not understand the content
- Use cases
 - Storing session data
 - User profiles and preferences
 - Shopping cart data



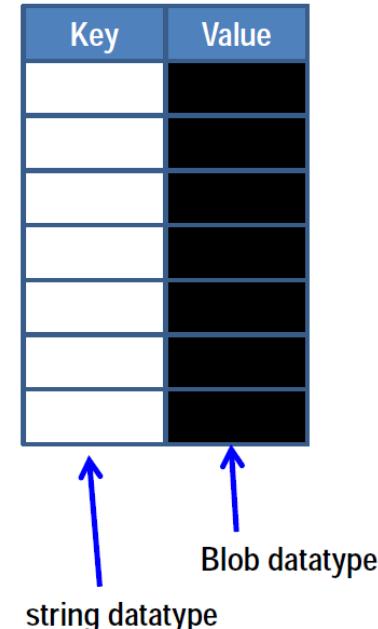
Key-value store

- **Pros:**

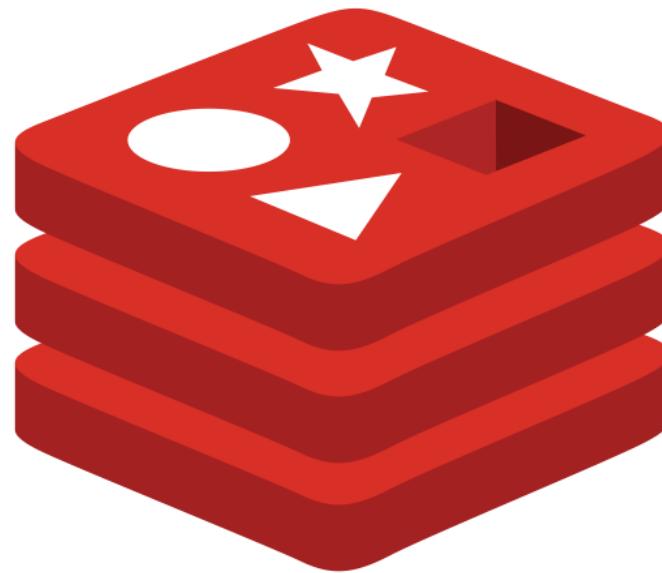
- Very fast
- Scalable
- Simple API
 - Put(key, value)
 - Get(key)
 - Delete(key)

- **Cons:**

- No way to query based on the content of the value



REDIS

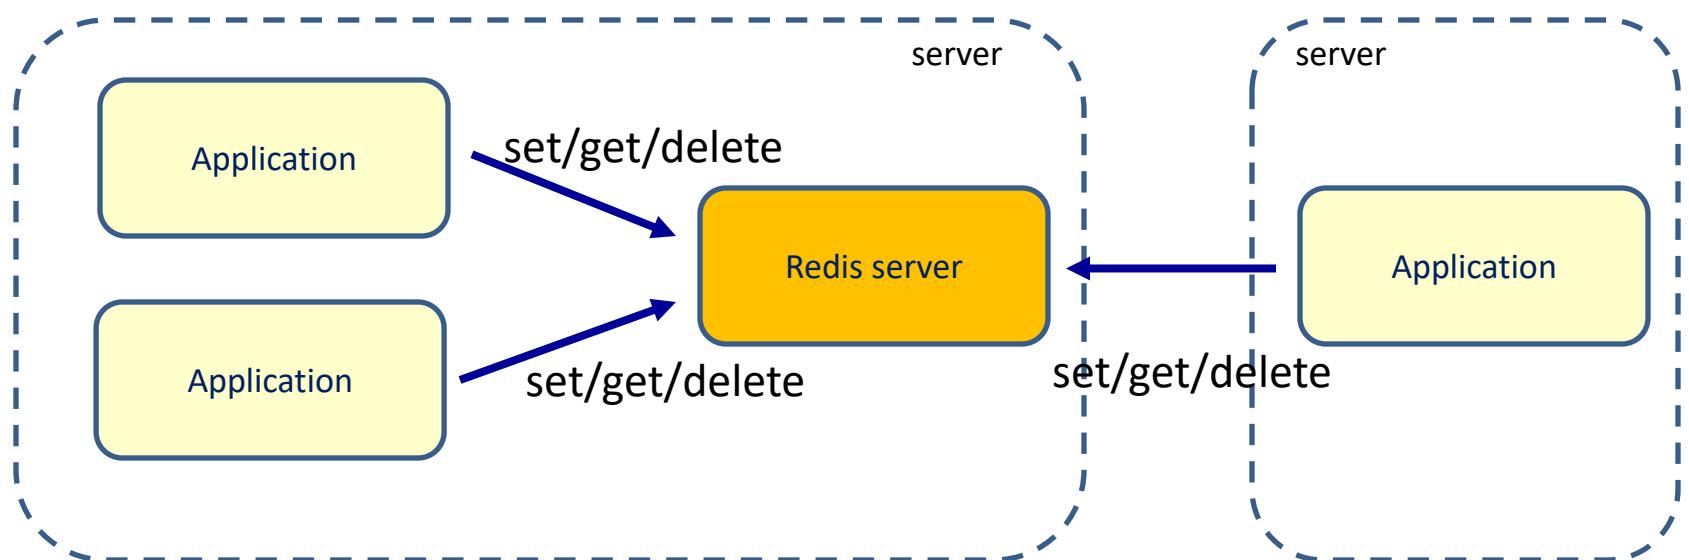


Redis

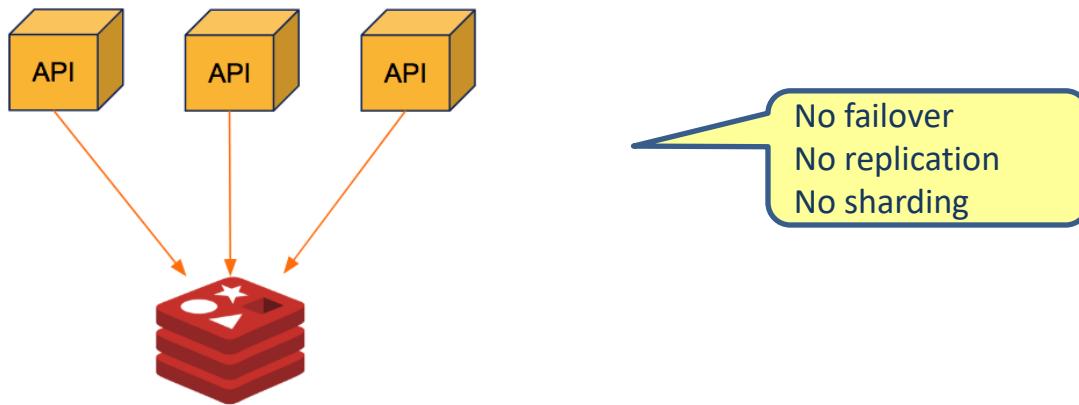
- REmote DIctionary Service
- Key-value store
- Very, very fast
 - Microseconds (not milliseconds)
- “in-memory” database
 - Can periodically write to disk
 - Size of data is limited by amount of memory in the sever



Using redis

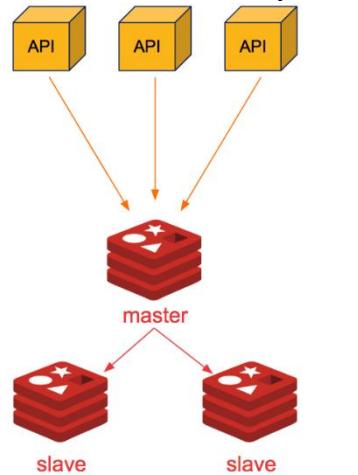


One redis instance



Master slave replication and sentinel

Master slave replication

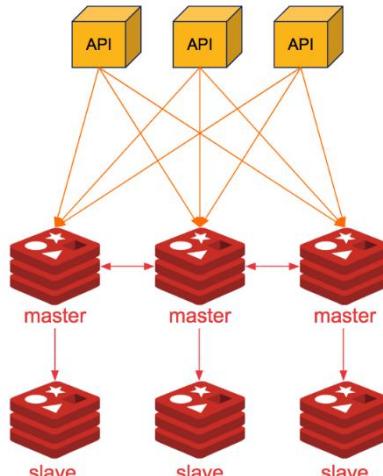


Failover because of replication

Slaves are read only

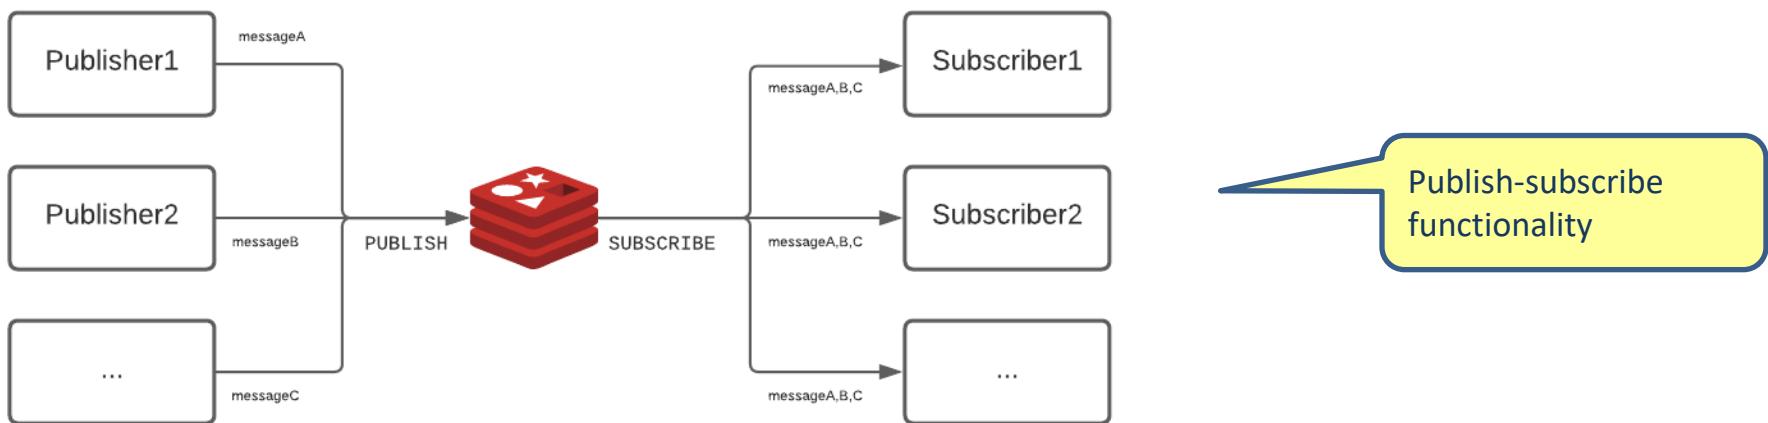
No automatic failover

Redis cluster



Automatic failover
Replication
Key based sharding

Redis streams



Redis advantages & disadvantages

- Advantages
 - Very fast in memory persistence
 - Scalable
 - High available
 - Message functionality
 - Pub-sub
- Disadvantages
 - Not for long term data persistence
 - You cannot query the value
 - No query language
 - Storage is limited to available memory



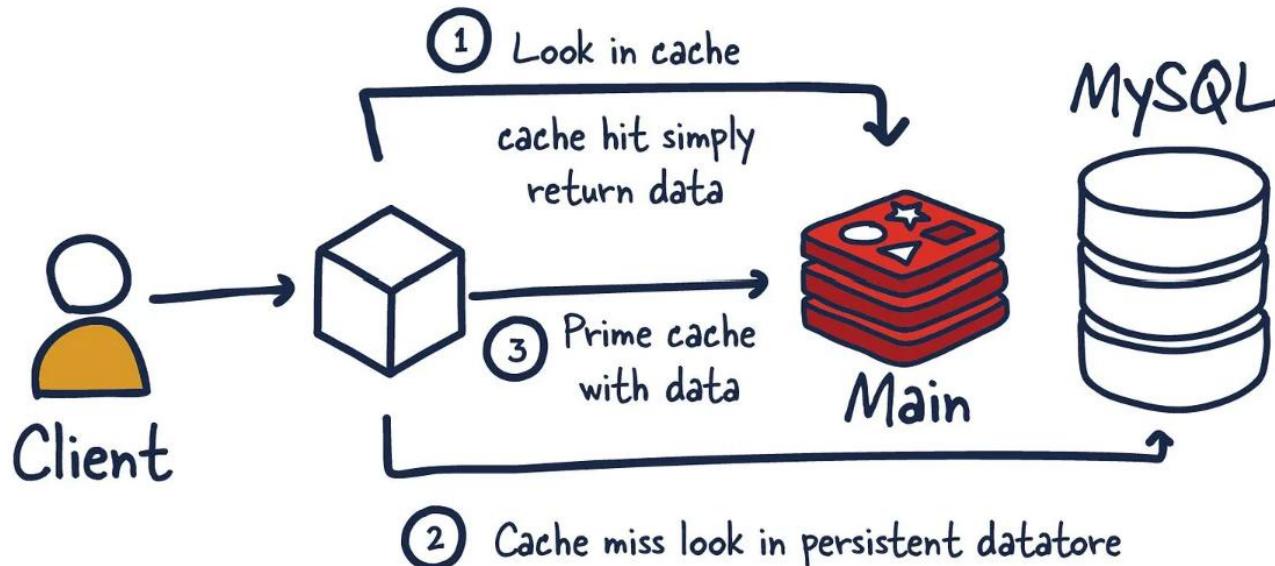
Redis use cases

- In memory cache
 - Data that changes infrequently and is requested often
 - Data that is less mission-critical and is frequently evolving
- Session management
- Queue systems
- Often used in combination with another database for permanent storage



Redis use cases

- Often used in combination with another database for permanent storage



SPRING BOOT REDIS EXAMPLE



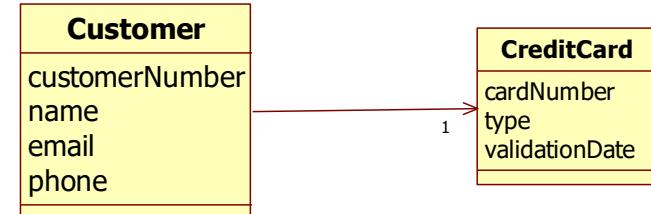
Domain classes

```
@RedisHash("Customer")
```

```
public class Customer {  
    @Id  
    private int customerNumber;  
    private String name;  
    private String email;  
    private String phone;  
    private CreditCard creditCard;
```

```
...  
  
public class CreditCard {  
    private String cardNumber;  
    private String type;  
    private String validationDate;
```

```
...
```



```
Application
```

```
main()  
run()
```

```
CustomerRepository
```

```
save()  
findById()  
findAll()
```

```
DI
```

Repository



```
@Repository  
public interface CustomerRepository extends CrudRepository<Customer, Integer> {}
```



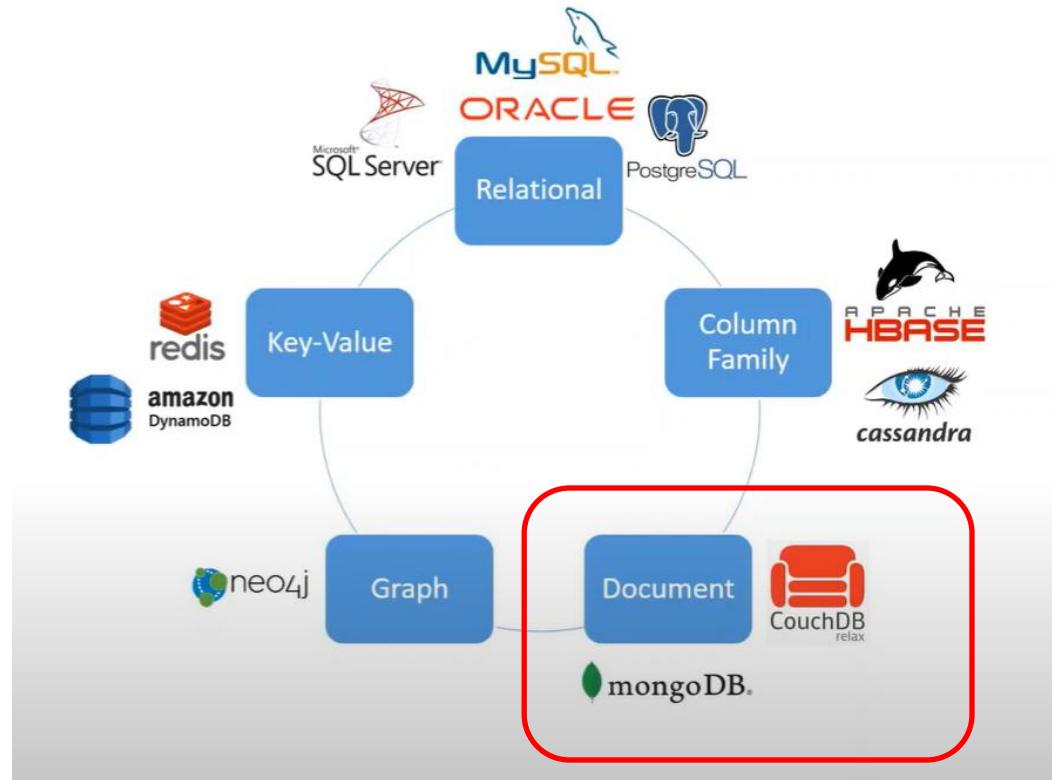
application.properties

```
# Redis configuration.  
spring.redis.host=localhost  
spring.redis.port=6379  
  
logging.level.root=ERROR  
logging.level.org.springframework=ERROR
```



Application

```
public void run(String... args) throws Exception {
    // create customer
    Customer customer = new Customer(101, "John doe", "johnd@acme.com", "0622341678");
    CreditCard creditCard = new CreditCard("12324564321", "Visa", "11/23");
    customer.setCreditCard(creditCard);
    customerRepository.save(customer);
    customer = new Customer(66, "James Johnson", "jj123@acme.com", "068633452");
    creditCard = new CreditCard("99876549876", "MasterCard", "01/24");
    customer.setCreditCard(creditCard);
    customerRepository.save(customer);
    //get customers
    System.out.println(customerRepository.findById(66).get());
    System.out.println(customerRepository.findById(101).get());
    System.out.println("-----All customers -----");
    System.out.println(customerRepository.findAll());
    //update customer
    customer = customerRepository.findById(101).get();
    customer.setEmail("jd@gmail.com");
    customerRepository.save(customer);
    //delete customer
    customerRepository.deleteById(66);
    System.out.println("-----All customers -----");
    System.out.println(customerRepository.findAll());
}
```



DOCUMENT DATABASE

Document database

- Store and retrieve documents
 - Like files in directories
- No schema
- Store documents in JSON



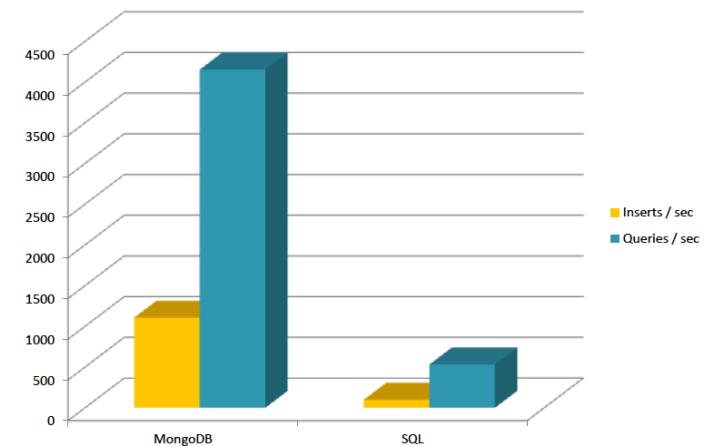


MONGODB



MongoDB

-
- Document database
 - Fast
 - Can handle large datasets



Mongo terminology

RDBMS	MongoDB
Database	Database
Table	Collection
Index	Index
Row	Document
Join	Embedding & Linking



Document data model (JSON)

Relational - Tables

Customer ID	First Name	Last Name	City
0	John	Doe	New York
1	Mark	Smith	San Francisco
2	Jay	Black	Newark
3	Meagan	White	London
4	Edward	Daniels	Boston

Account Number	Branch ID	Account Type	Customer ID
10	100	Checking	0
11	101	Savings	0
12	101	IRA	0
13	200	Checking	1
14	200	Savings	1
15	201	IRA	2

Document - Collections

```
{   customer_id : 1,  
    first_name : "Mark",  
    last_name : "Smith",  
    city : "San Francisco",  
    accounts : [    {  
        account_number : 13,  
        branch_ID : 200,  
        account_type : "Checking"  
    },  
    {   account_number : 14,  
        branch_ID : 200,  
        account_type : "IRA",  
        beneficiaries: [...]  
    } ]  
}
```



Documents are rich data structures

```
{  
    first_name: 'Paul',  
    surname: 'Miller',  
    cell: 447557505611,  
    city: 'London',  
    location: [45.123,47.232],  
    Profession: ['banking', 'finance',  
    'trader'],  
    cars: [  
        { model: 'Bentley',  
         year: 1973,  
         value: 100000, ... },  
        { model: 'Rolls Royce',  
         year: 1965,  
         value: 330000, ... }  
    ]  
}
```

Fields

String
Number
Geo-Coordinates

Typed field values

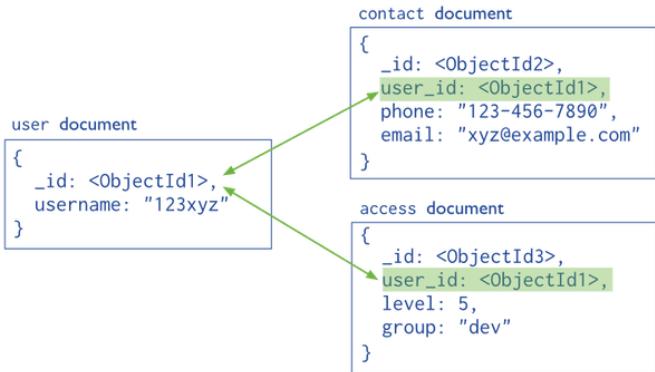
Fields can contain arrays

Fields can contain an array of sub-documents



References vs. embedded data

Reference



Embedded data



Schema free

```
{name: "will",  
eyes: "blue",  
birthplace: "NY",  
aliases: ["bill", "la  
ciacco"],  
gender: "????",  
boss:"ben"}
```

```
{name: "jeff",  
eyes: "blue",  
height: 72,  
boss: "ben"}
```

```
{name: "ben",  
hat:"yes"}
```

```
{name: "brendan",  
aliases: ["el diablo"]}
```

```
{name: "matt",  
pizza: "DiGiorno",  
height: 72,  
boss: 555-555-1212}
```



Find() method

SQL SELECT Statements

`SELECT * FROM users`

`SELECT id, user_id, status FROM users`

`SELECT user_id, status FROM users`

`SELECT * FROM users WHERE status = "A"`

`SELECT user_id, status FROM users WHERE status = "A"`

`SELECT * FROM users WHERE status != "A"`

`SELECT * FROM users WHERE status = "A" AND age = 50`

`SELECT * FROM users WHERE status = "A" OR age = 50`

`SELECT * FROM users WHERE age > 25`

MongoDB find() Statements

`db.users.find()`

`db.users.find({}, { user_id: 1, status: 1 })`

`db.users.find({}, { user_id: 1, status: 1, _id: 0 })`

`db.users.find({ status: "A" })`

`db.users.find({ status: "A" }, { user_id: 1, status: 1, _id: 0 })`

`db.users.find({ status: { $ne: "A" } })`

`db.users.find({ status: "A", age: 50 })`

`db.users.find({ $or: [{ status: "A" } , { age: 50 }] })`

`db.users.find({ age: { $gt: 25 } })`



How to structure data in mongo?

To get high performance,
every query should be
done on 1 node only



Data is sharded
over multiple nodes

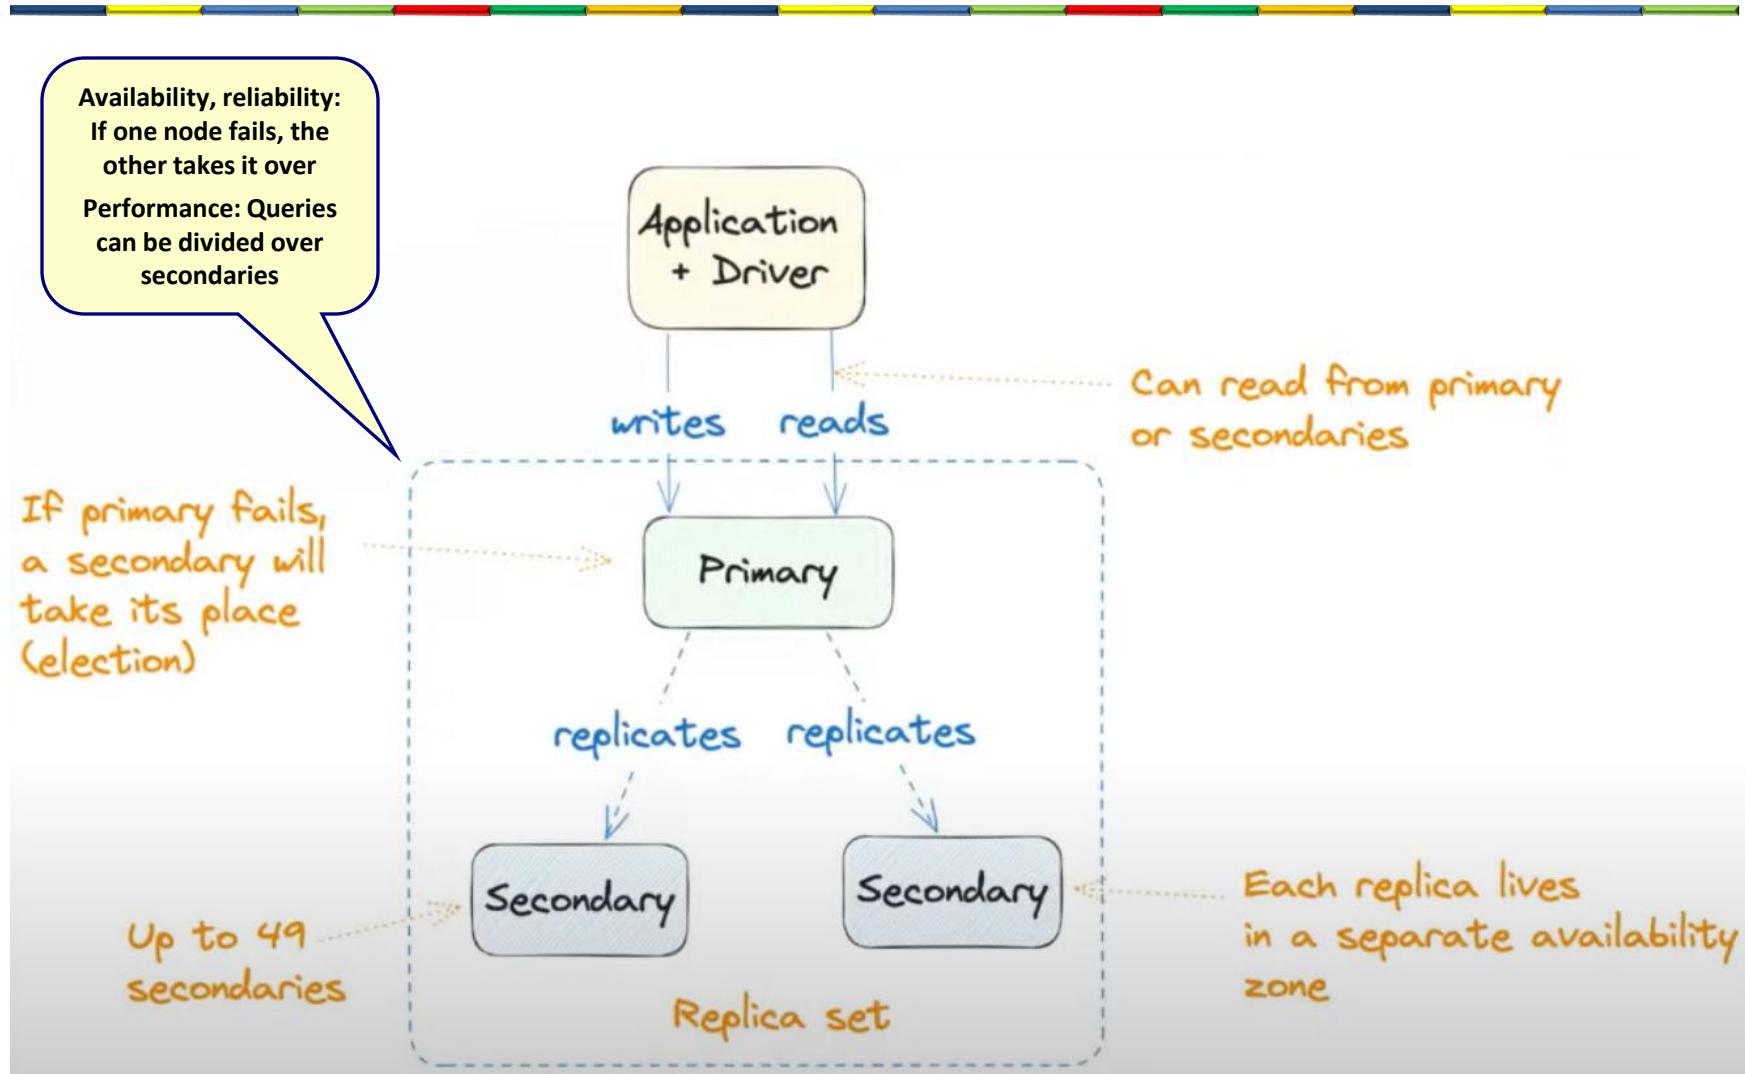


Structure the data
according your
queries

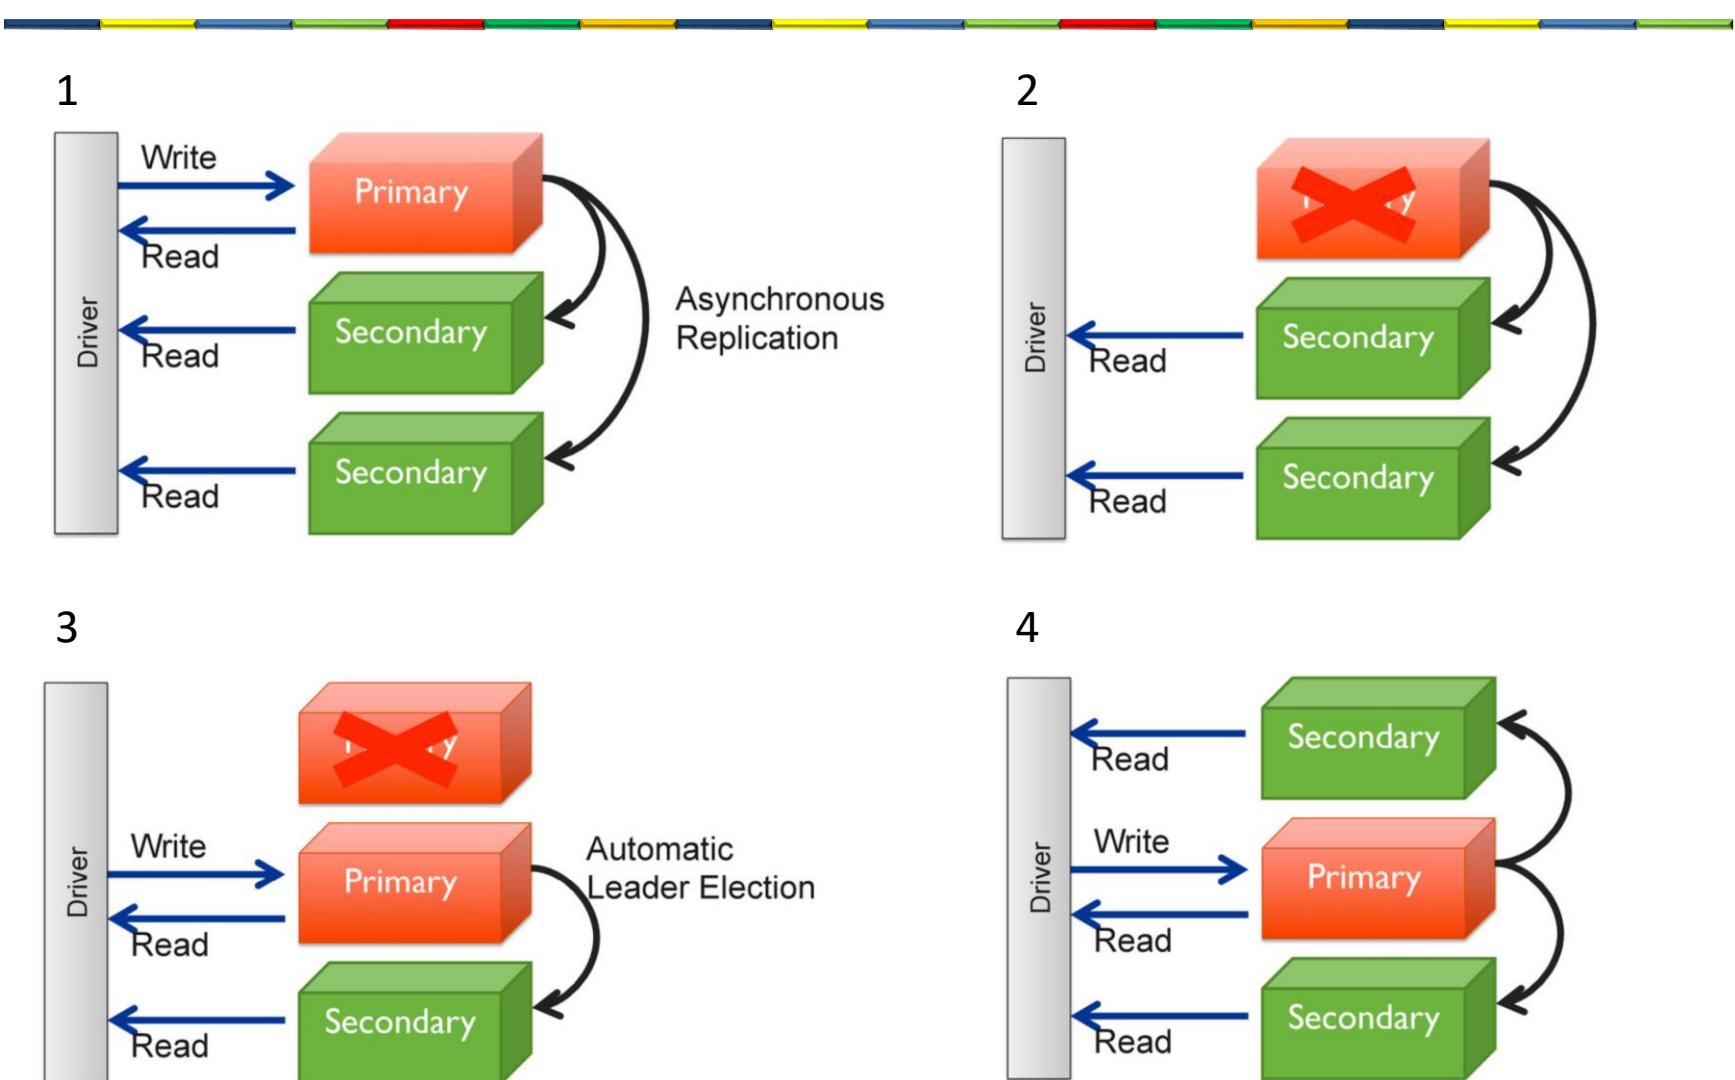
Consequence: data
duplication



MongoDB replication

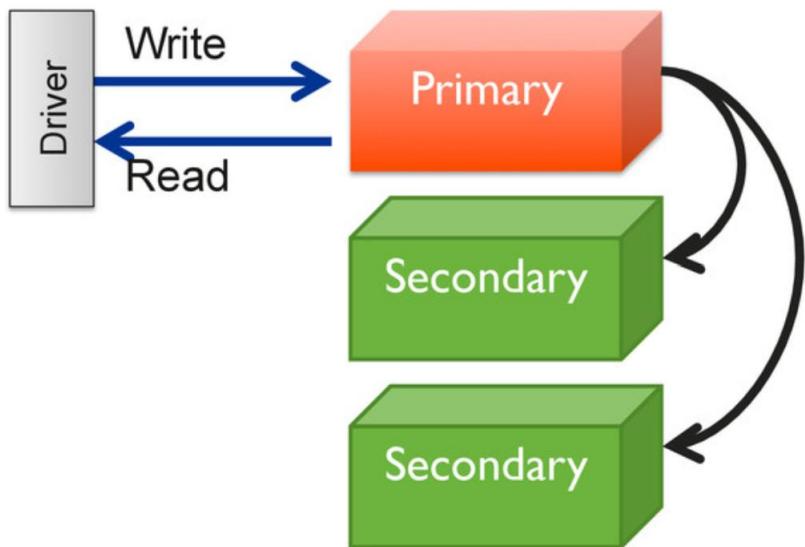


Replica sets

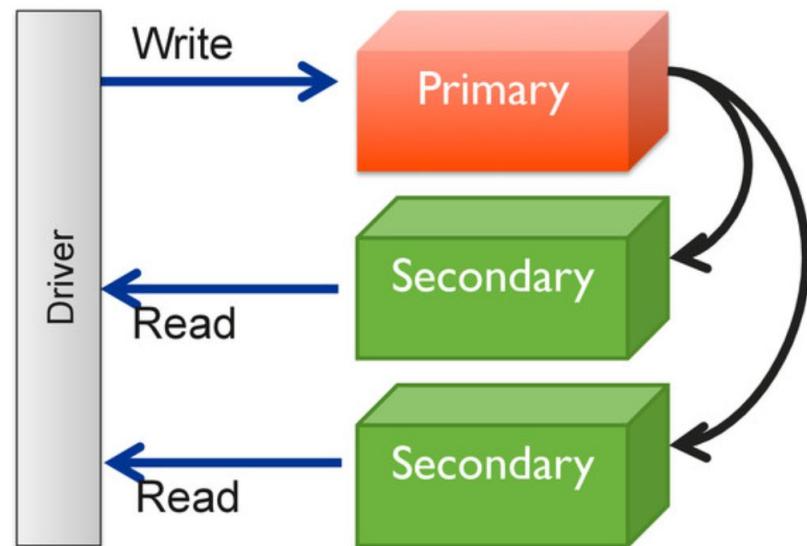


Consistency

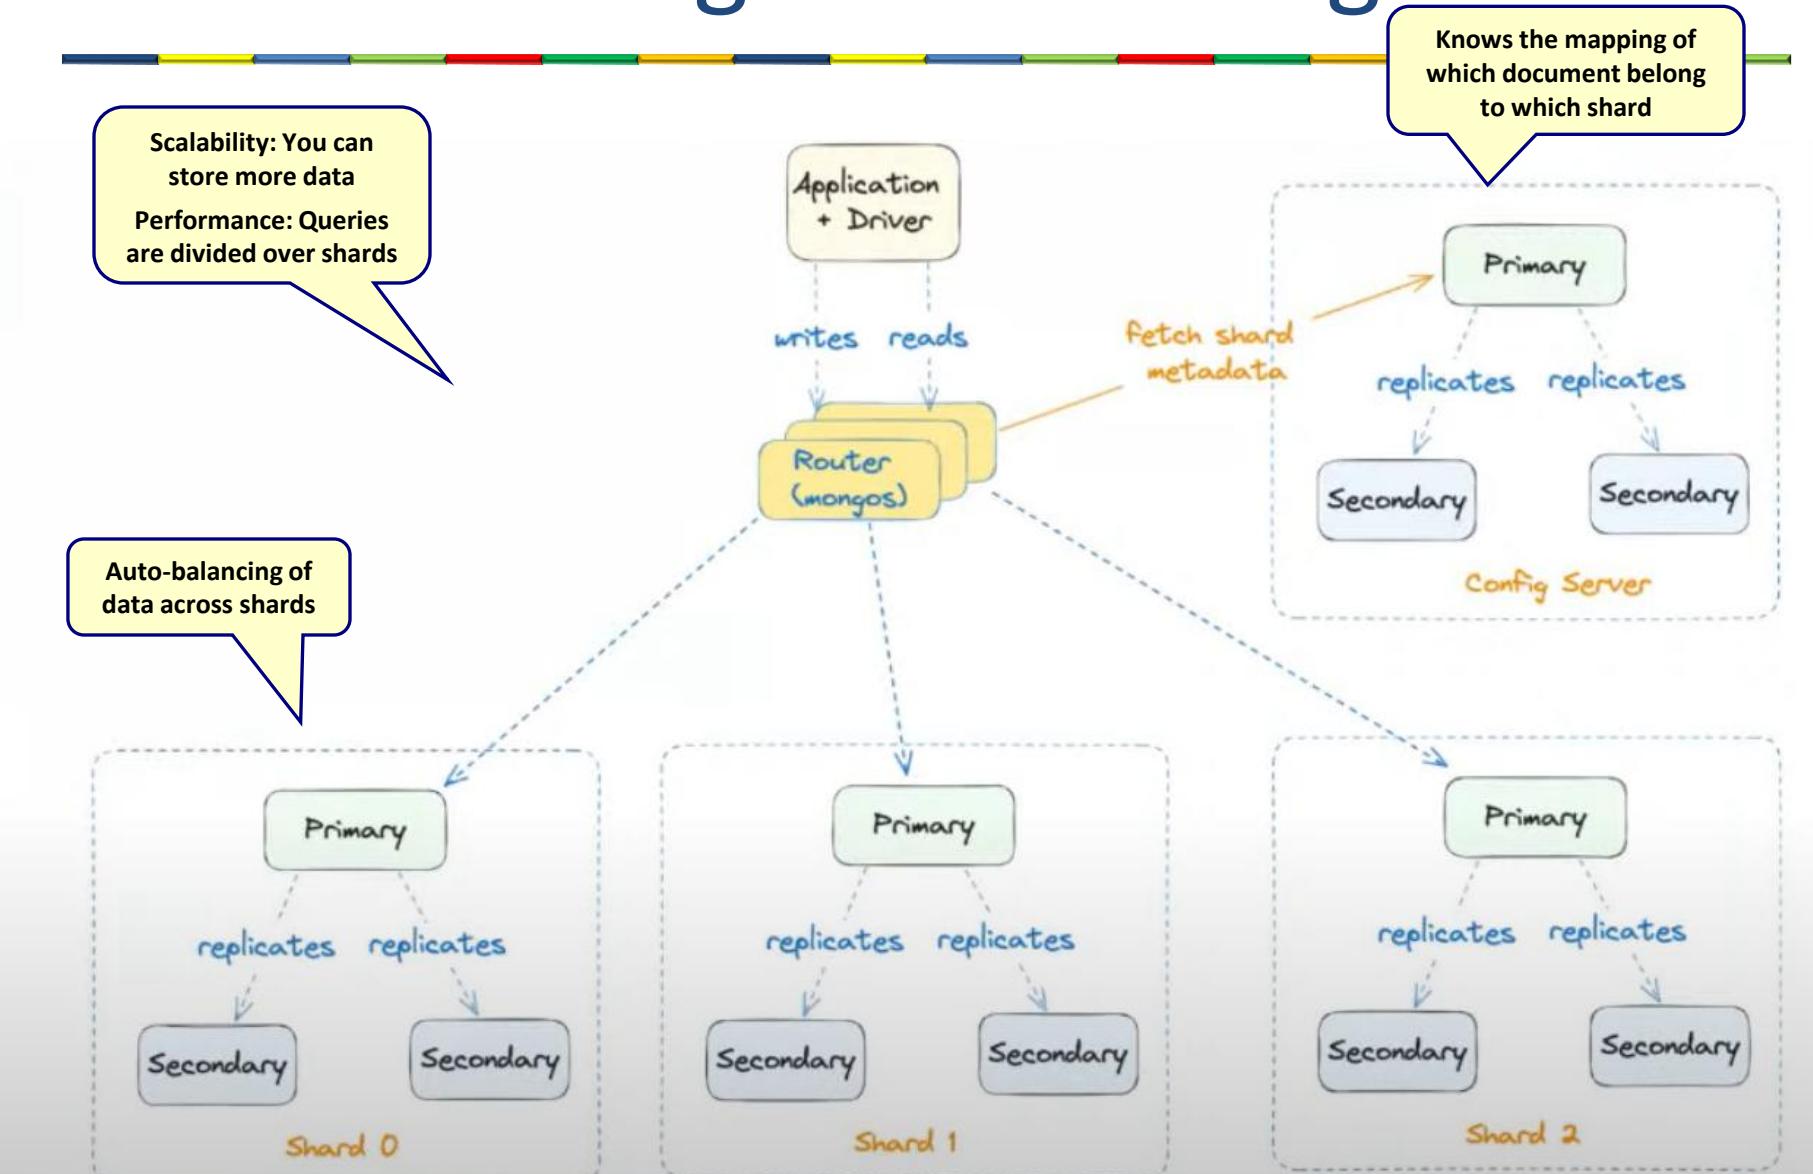
Strong consistency



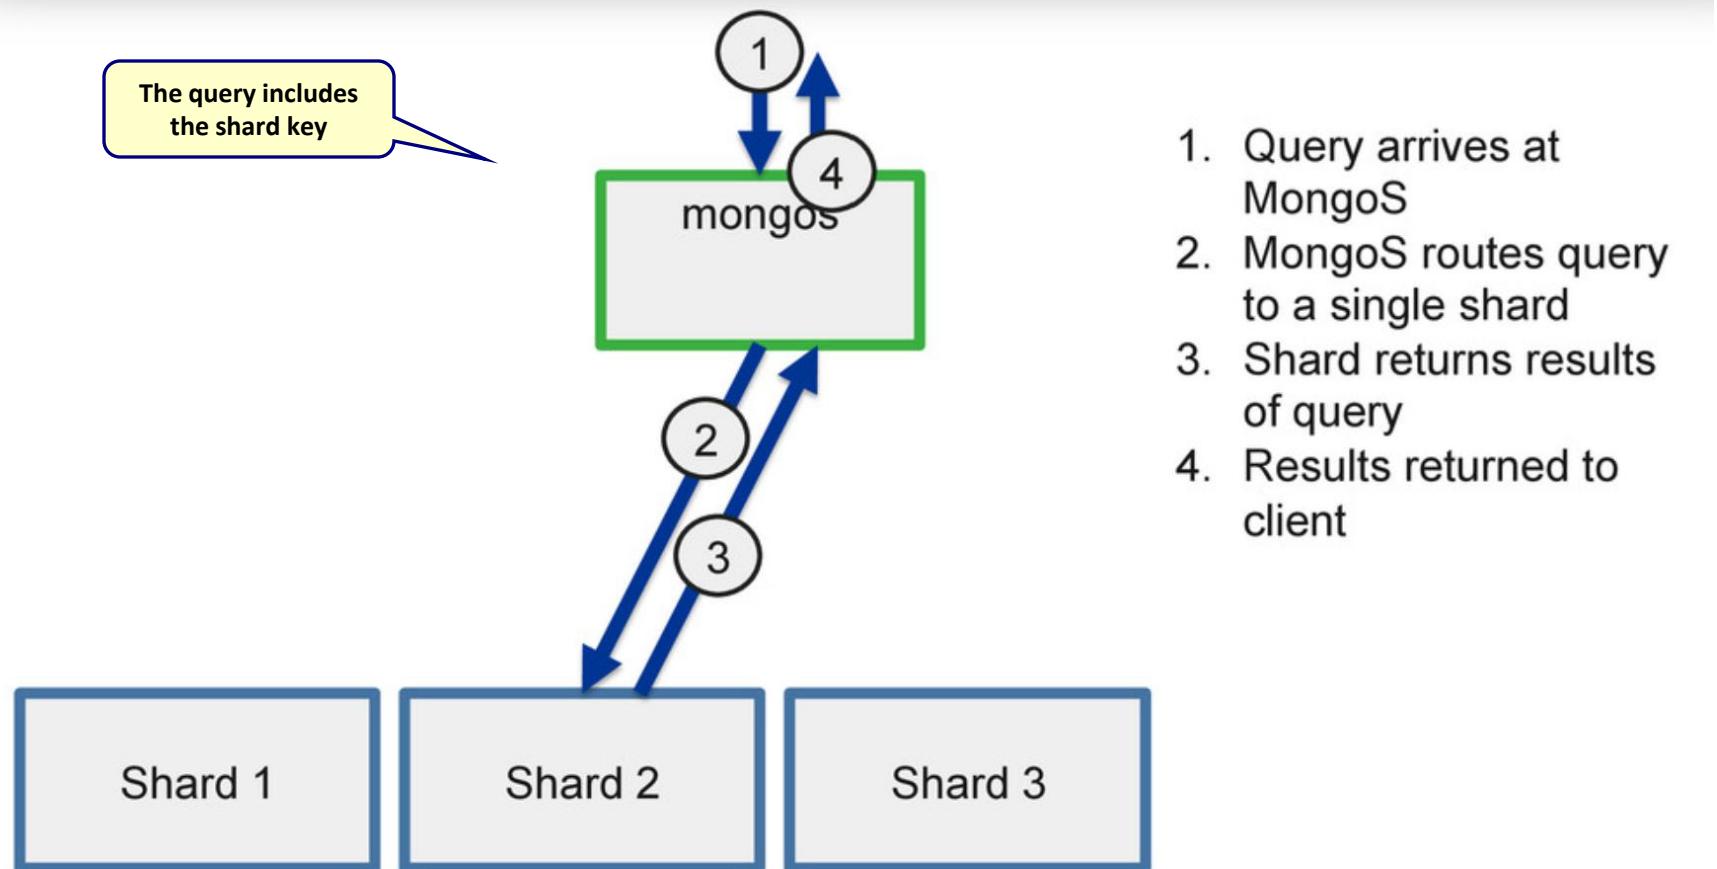
Eventual consistency



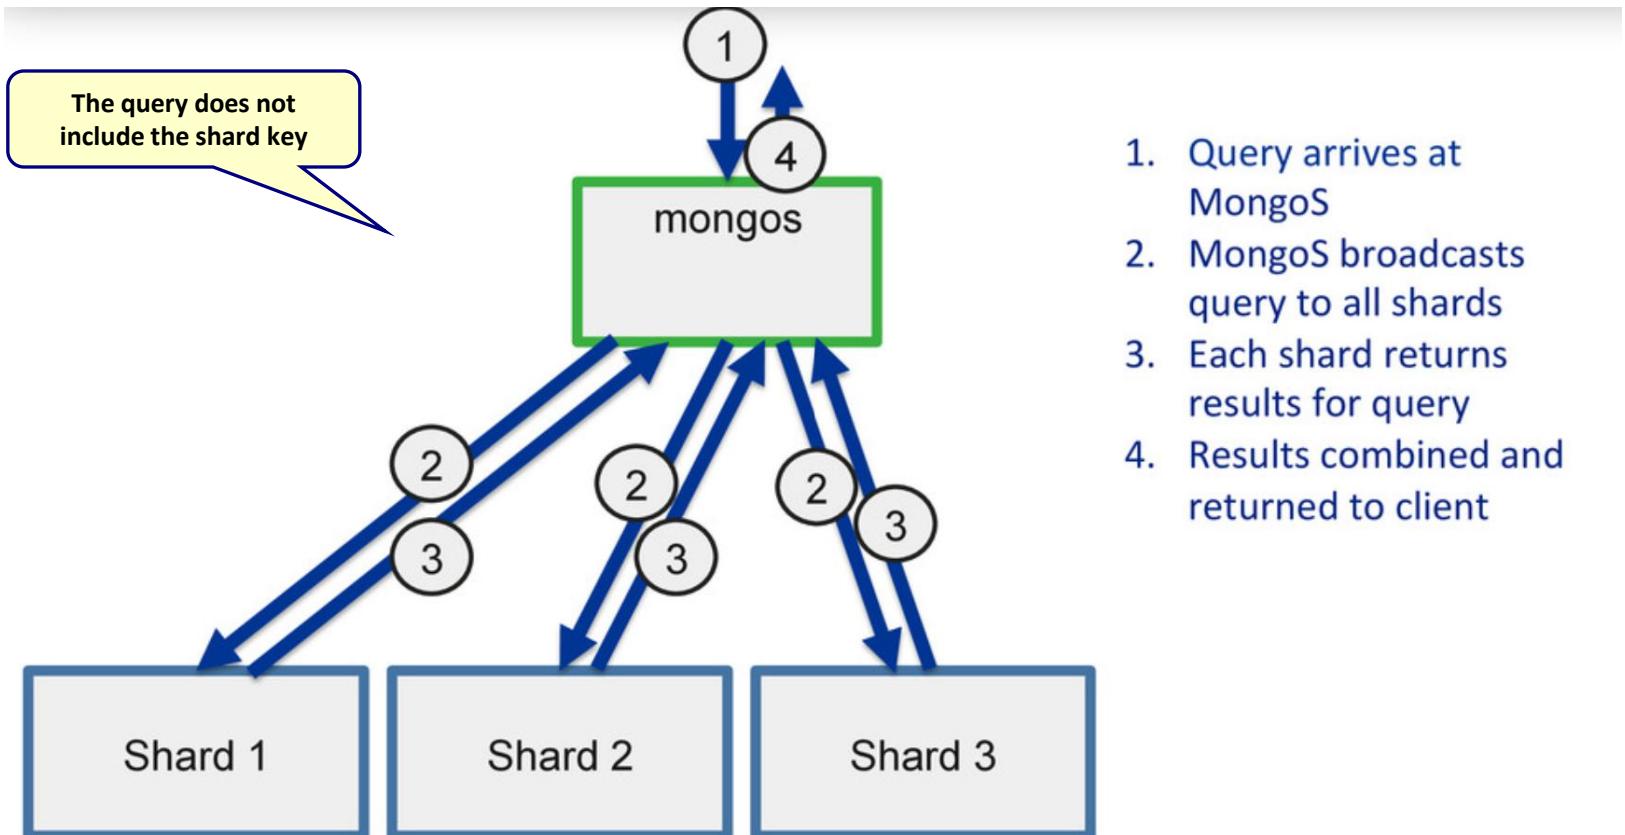
MongoDB sharding



Routed query



Scatter gather



Mongo advantages & disadvantages

- Advantages
 - Very fast
 - Scalable
 - High available
 - Handles unstructured data
 - Rich query language
 - Indexing
 - No schema
 - ACID transaction support
- Disadvantages
 - Limited joins
 - Data redundancy
 - Document size limit (16 MB)
 - Limited nested documents levels (100)



Mongo use cases

- IoT
 - Large growing data sets, flexible schema
- Real-time analytics
 - Large growing data sets, flexible schema
- Content Management
 - Large growing data sets, flexible schema
- Product catalog
 - Large growing data sets, flexible schema



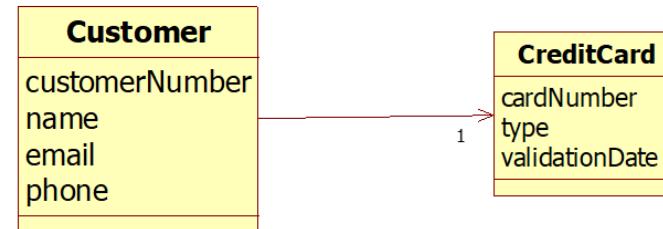
SPRING BOOT MONGO EXAMPLE



Customer

```
@Document  
public class Customer {  
    @Id  
    private int customerNumber;  
    private String name;  
    private String email;  
    private String phone;  
    private CreditCard creditCard;  
    ...}
```

```
public class CreditCard {  
    private String cardNumber;  
    private String type;  
    private String validationDate;  
    ...}
```



Repository

```
@Repository
public interface CustomerRepository extends MongoRepository<Customer,
Integer> {
    Customer findByPhone(String phone);
    Customer findByName(String name);

    @Query("{email : ?0}")
    Customer findCustomerWithPhone(String email);

    @Query("{creditCard.type : ?0}")
    List<Customer> findCustomerWithCreditCardType(String ctype);
}
```



Application (1/2)

```
public void run(String... args) throws Exception {
    //create customer
    Customer customer = new Customer(101,"John doe", "johnd@acme.com", "0622341678");
    CreditCard creditCard = new CreditCard("12324564321", "Visa", "11/23");
    customer.setCreditCard(creditCard);
    customerRepository.save(customer);
    customer = new Customer(109,"John Jones", "jones@acme.com", "0624321234");
    creditCard = new CreditCard("657483342", "Visa", "09/23");
    customer.setCreditCard(creditCard);
    customerRepository.save(customer);
    customer = new Customer(66,"James Johnson", "jj123@acme.com", "068633452");
    creditCard = new CreditCard("99876549876", "MasterCard", "01/24");
    customer.setCreditCard(creditCard);
    customerRepository.save(customer);
    //get customers
    System.out.println(customerRepository.findById(66).get());
    System.out.println(customerRepository.findById(101).get());
```

Application (2/2)

```
System.out.println("-----All customers -----");
System.out.println(customerRepository.findAll());
//update customer
customer = customerRepository.findById(101).get();
customer.setEmail("jd@gmail.com");
customerRepository.save(customer);
System.out.println("-----find by phone -----");
System.out.println(customerRepository.findByPhone("0622341678"));
System.out.println("-----find by email -----");
System.out.println(customerRepository.findCustomerWithPhone("jj123@acme.com"));
System.out.println("-----find customers with a certain type of creditcard -----");
List<Customer> customers = customerRepository.findCustomerWithCreditCardType("Visa");
for (Customer cust : customers){
    System.out.println(cust);
}

}
```

Mongo collections

The screenshot shows the MongoDB Compass interface connected to the 'testdb.customer' database. The left sidebar displays the 'Local' connection details, including the host 'localhost:27017', cluster 'Standalone', and edition 'MongoDB 3.2.19-14-ge59d00a Community'. The 'testdb' database is selected, and the 'customer' collection is currently viewed. The right panel shows the 'Documents' tab for the 'testdb.customer' collection, displaying three document entries. Each document includes fields such as _id, name, email, phone, creditCard (an object containing cardNumber, type, validationDate, and _class), and a timestamp.

MongoDB Compass - localhost:27017/testdb.customer

Connect View Collection Help

Local

3 DBS 12 COLLECTIONS

HOST
localhost:27017

CLUSTER
Standalone

EDITION
MongoDB 3.2.19-14-
ge59d00a Community

Filter your data

local
test
testdb
customer

testdb.customer

Documents Aggregations

FILTER { field: 'value' }

ADD DATA VIEW

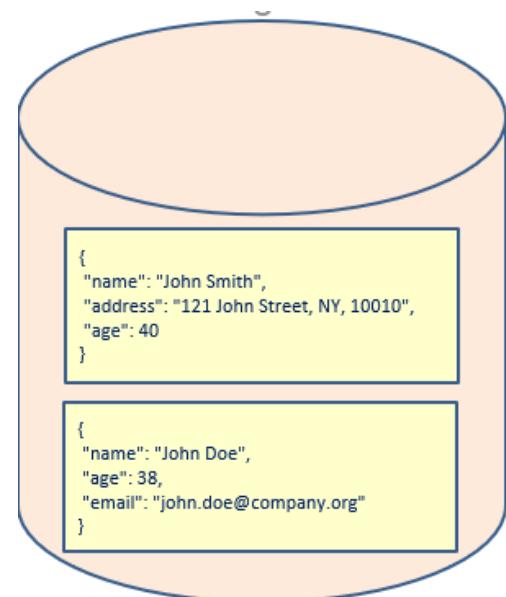
`_id: 101
name: "John Doe"
email: "j@gmail.com"
phone: "0622341678"
creditCard: Object
 cardNumber: "12324564321"
 type: "Visa"
 validationDate: "11/23"
 _class: "customers.domain.Customer"`

`_id: 109
name: "John Jones"
email: "jones@acme.com"
phone: "0624321234"
creditCard: Object
 cardNumber: "657483342"
 type: "Visa"
 validationDate: "09/23"
 _class: "customers.domain.Customer"`

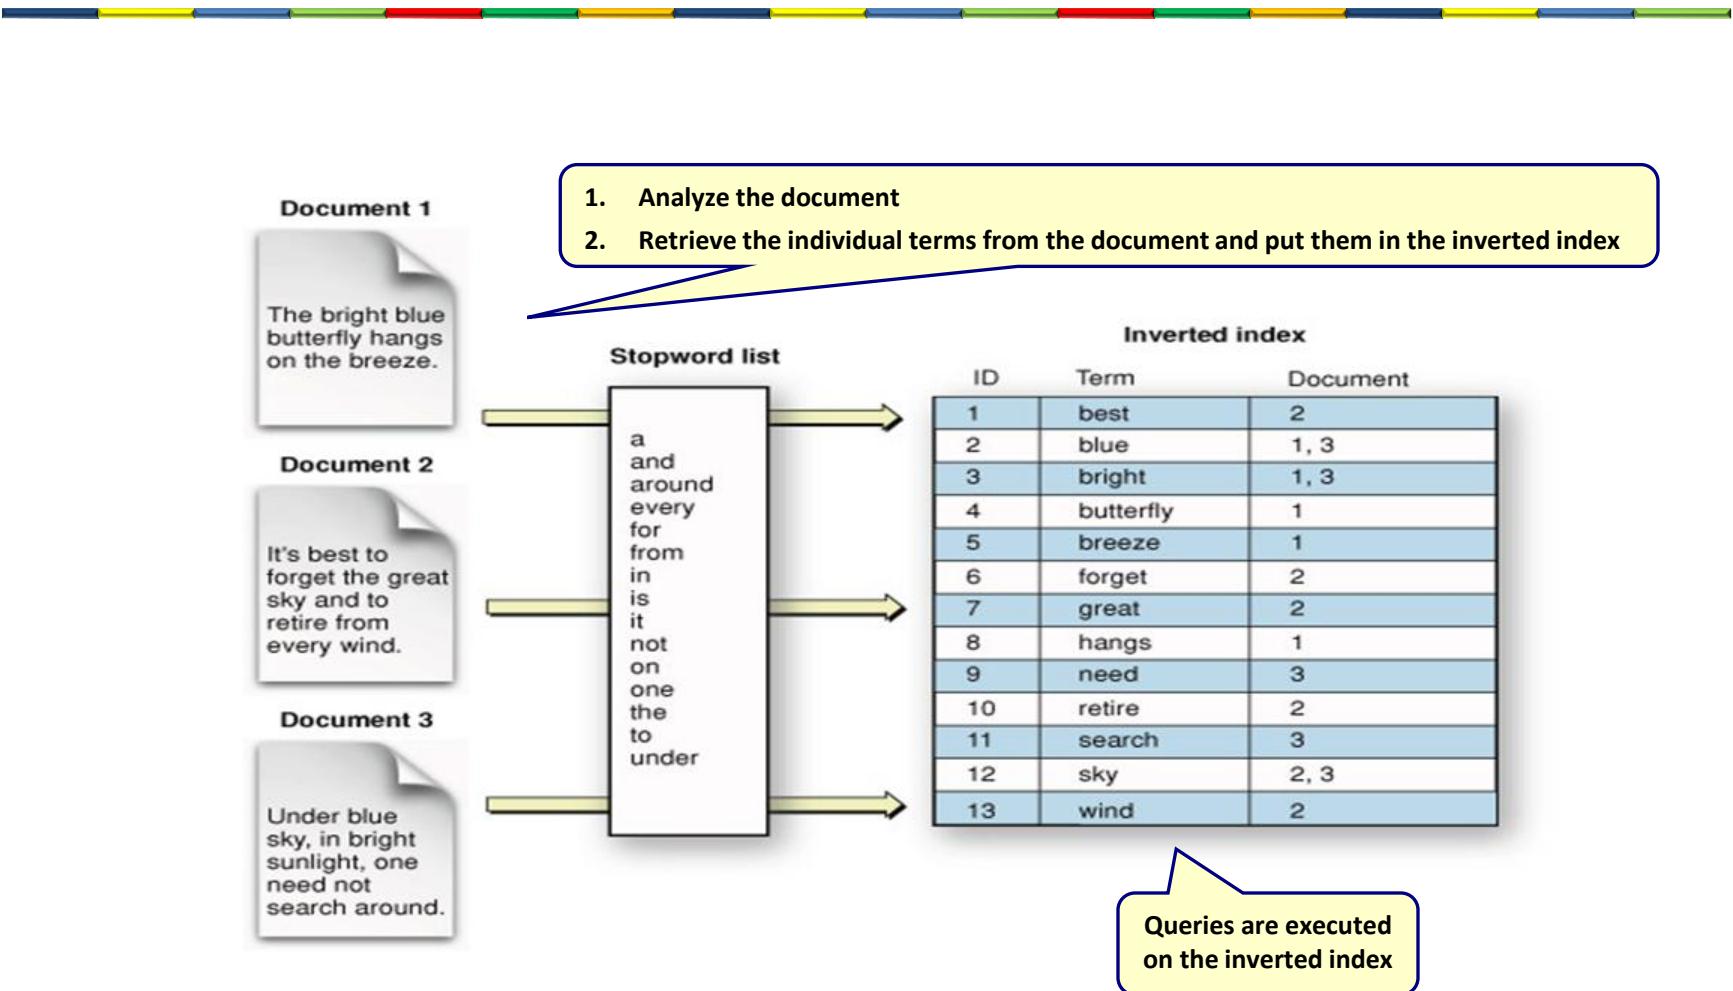
`_id: 66
name: "James Johnson"
email: "jj123@acme.com"
phone: "068633452"
creditCard: Object
 cardNumber: "99876549876"
 type: "MasterCard"
 validationDate: "01/24"
 _class: "customers.domain.Customer"`

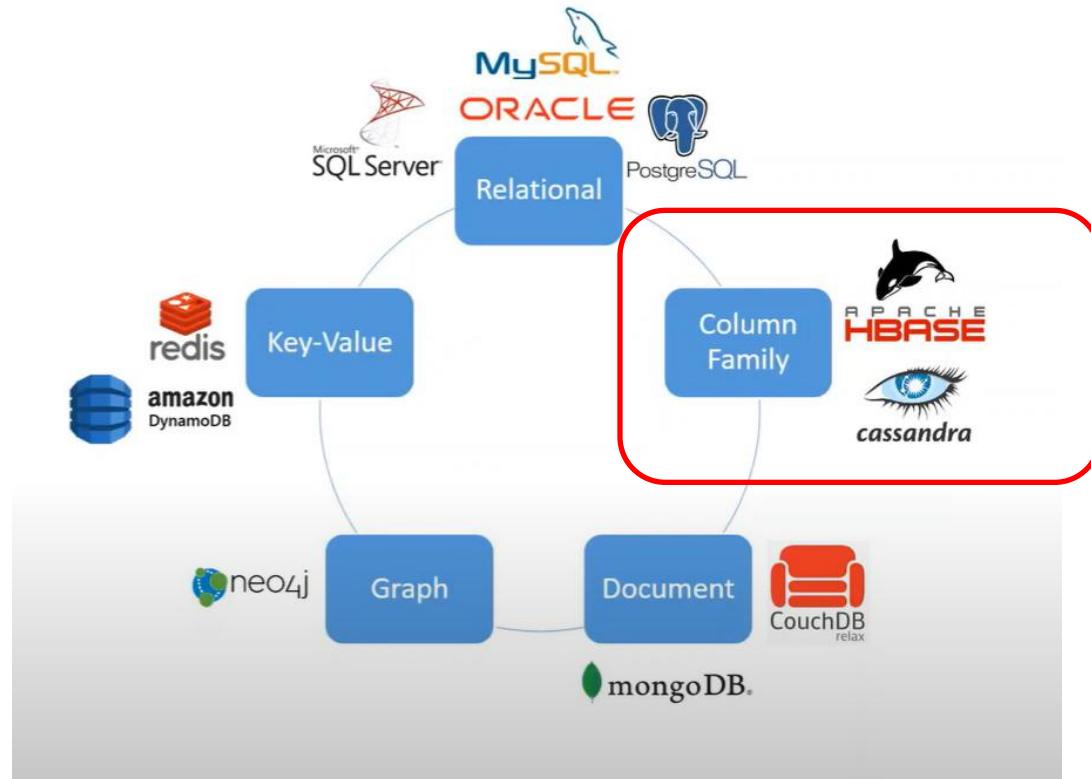
Elasticsearch

- Database
 - Data is stored as documents
 - Data is structured in JSON format
- Full text search engine



Inverted index

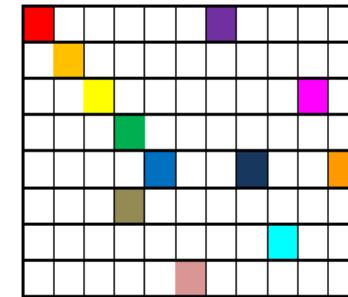




COLUMN FAMILY DATABASE

Column family

- Place data in a certain column
- Ideal for high-variability data sets
- Column families allow to query all columns that have a specific property or properties
- Allow new columns to be inserted without doing an "alter table"

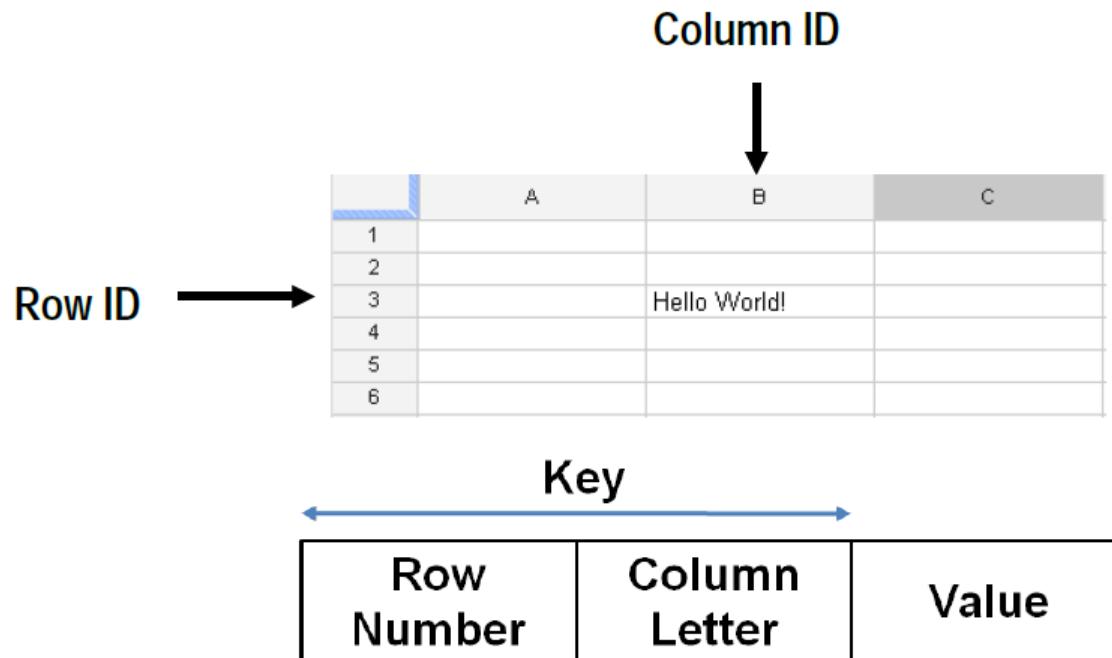


- Players
- Cassandra
- HBase
- Google BigTable

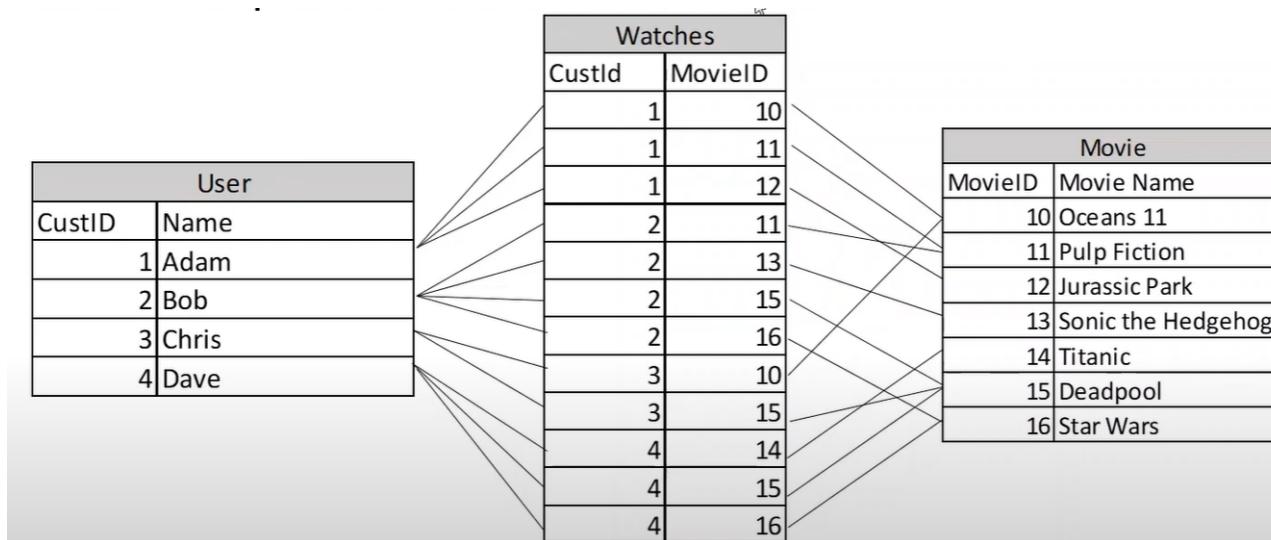


Excel

- The key is the row number and the column letter



Relational



What movies has Adam watched?



Need to join 3 tables (slow)

Who watched Deadpool?



Need to join 3 tables (slow)



Column family

User	Movies			
	Oceans 11	Pulp Fiction	Jurassic Park	
Adam	TRUE	TRUE	TRUE	
	Pulp Fiction	Sonic	Deadpool	Star Wars
Bob	TRUE	TRUE	TRUE	TRUE
	Oceans 11	Deadpool		
Chris	TRUE	TRUE		
	Titanic	Deadpool	Star Wars	
Dave	TRUE	TRUE	TRUE	

Structure the data
according your
queries

What movies has Adam watched?



Very easy and fast

Who watched Deadpool?



Difficult



CASSANDRA



Cassandra

- Column family NoSQL database
 - Uses tables, primary keys, queries, ...
- Unlimited elastically scalable
- Always available (no downtime)
- High performance
- Fault tolerance



High performance

MySQL Comparision:

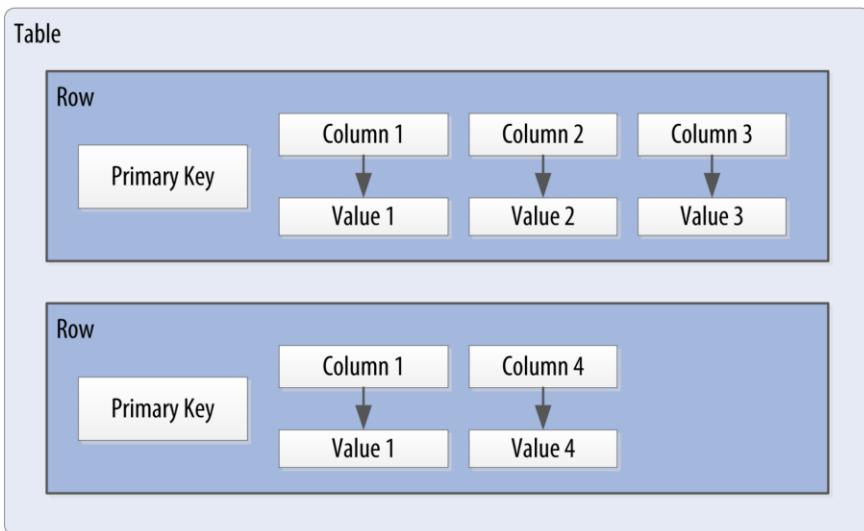


Statistics based on 50 GB Data

	Cassandra	MySQL
Average Write	0.12 ms	~300 ms
Average Read	15 ms	~350 ms



Cassandra data model



UserProfile

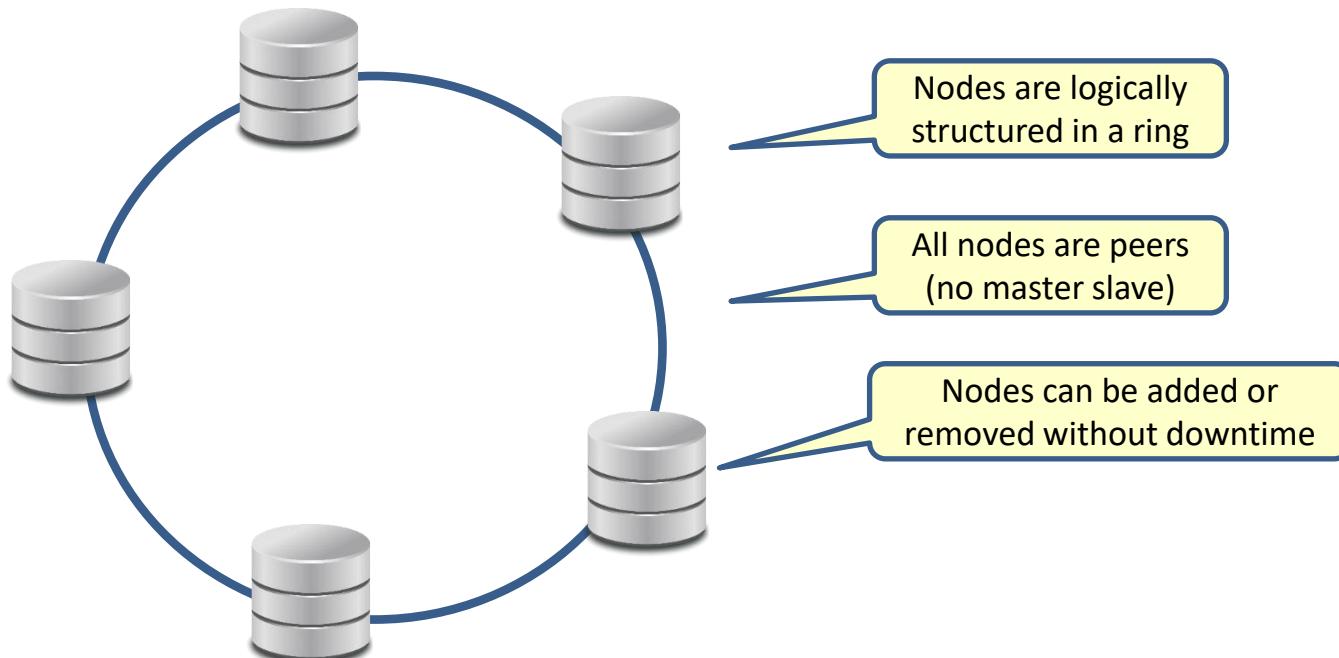
Bob	emailAddress	gender	age
	bob@example.com	male	35
	1465676582	1465676582	1465676582

Britney	emailAddress	gender
	brit@example.com	female
	1465676432	1465676432

Tori	emailAddress	country	hairColor
	tori@example.com	Sweden	Blue
	1435636158	1435636158	1465633654



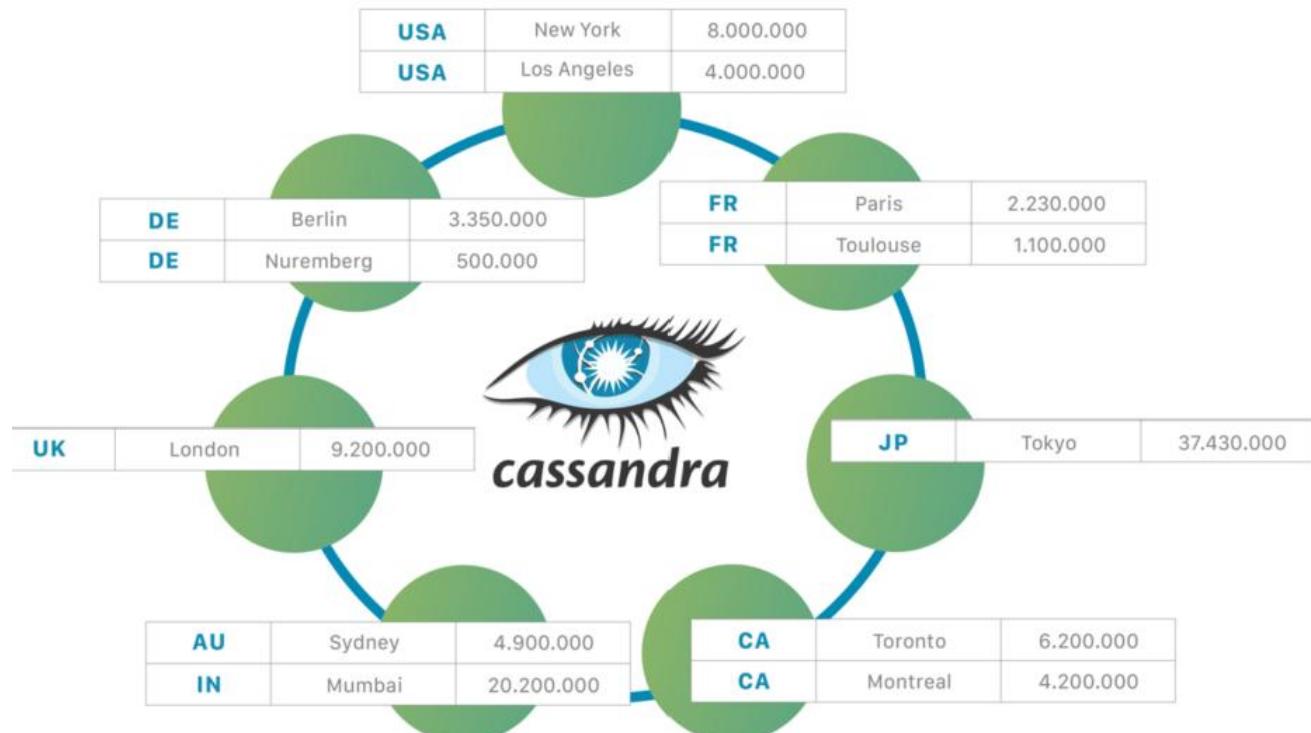
Cassandra cluster of nodes



Cassandra partitioning

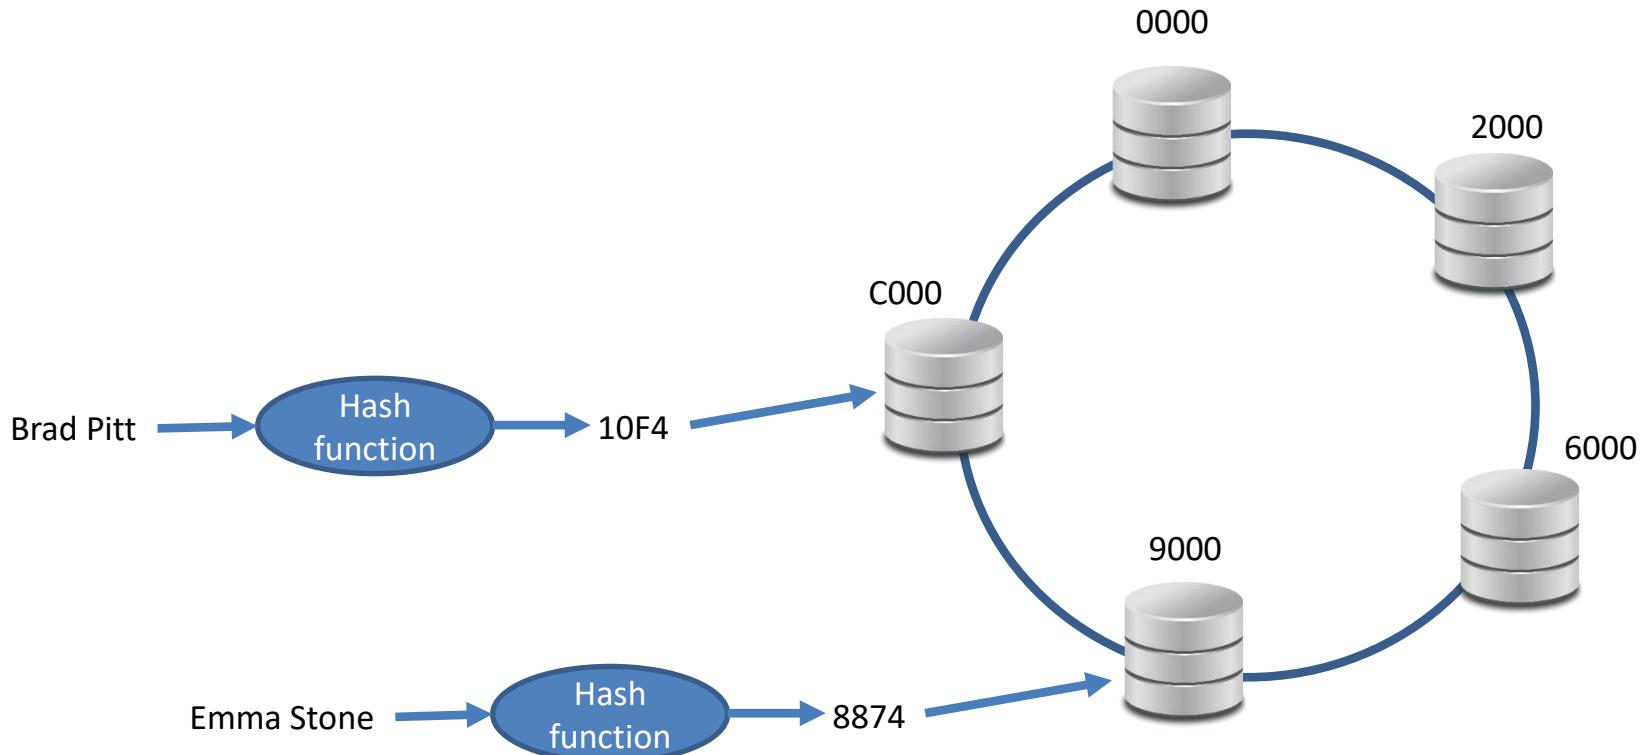
COUNTRY	CITY	POPULATION
USA	New York	8.000.000
USA	Los Angeles	4.000.000
FR	Paris	2.230.000
DE	Berlin	3.350.000
UK	London	9.200.000
AU	Sydney	4.900.000
DE	Nuremberg	500.000
CA	Toronto	6.200.000
CA	Montreal	4.200.000
FR	Toulouse	1.100.000
JP	Tokyo	37.430.000
IN	Mumbai	20.200.000

Partition Key

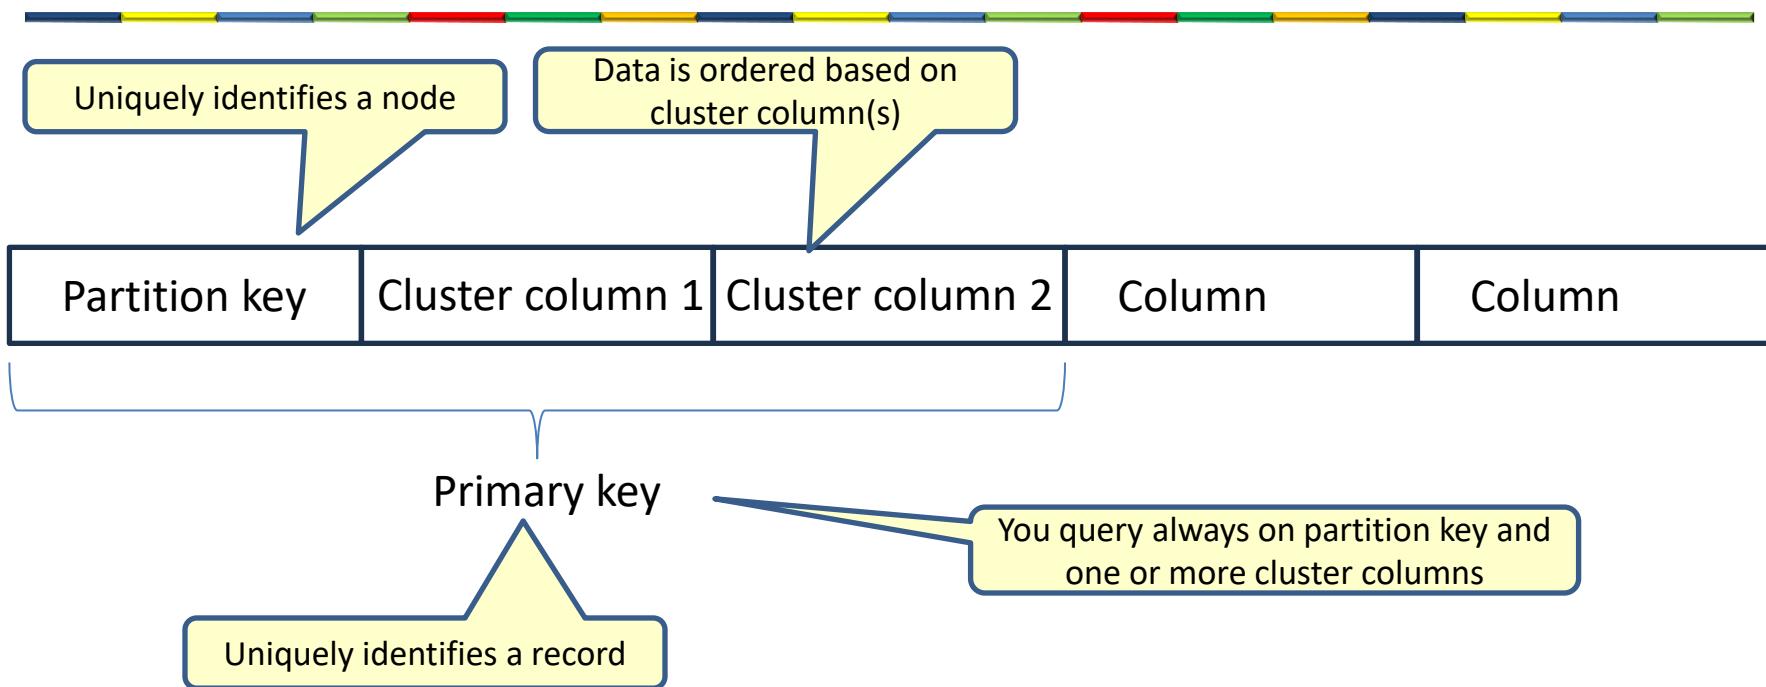


Partition key

- Decides where to place this record



Primary key



```
CREATE TABLE movies_by_actor (
    first_name TEXT,
    last_name TEXT,
    title TEXT,
    year INT,
    PRIMARY KEY ((first_name, last_name), title, year)
); ↵
```

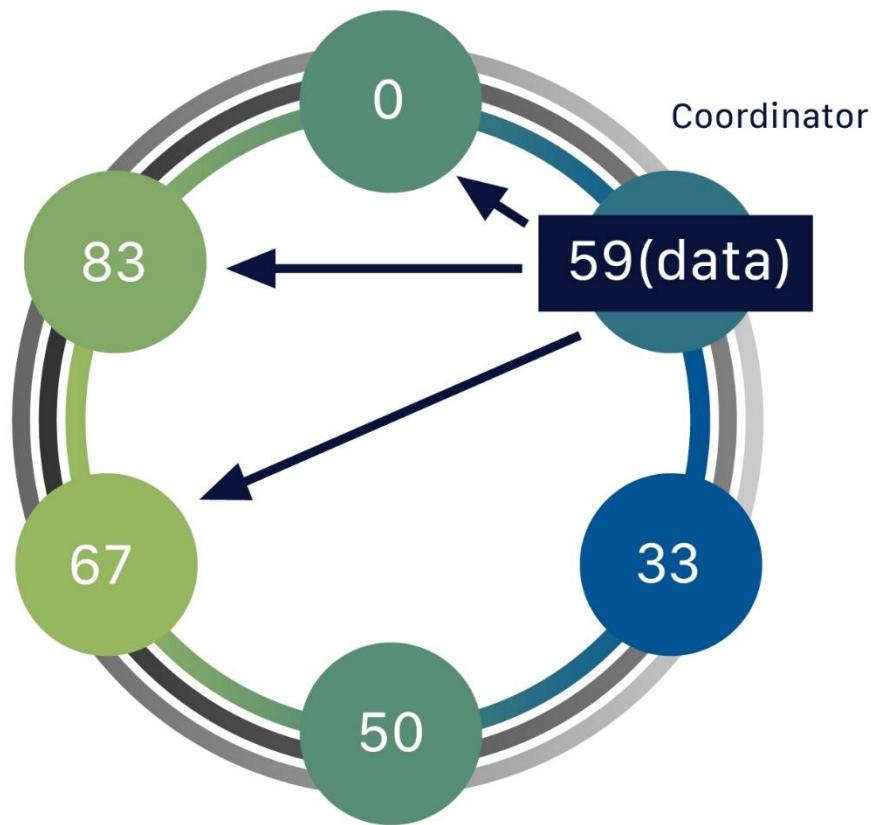
partition
key(s)

clustering
key(s)

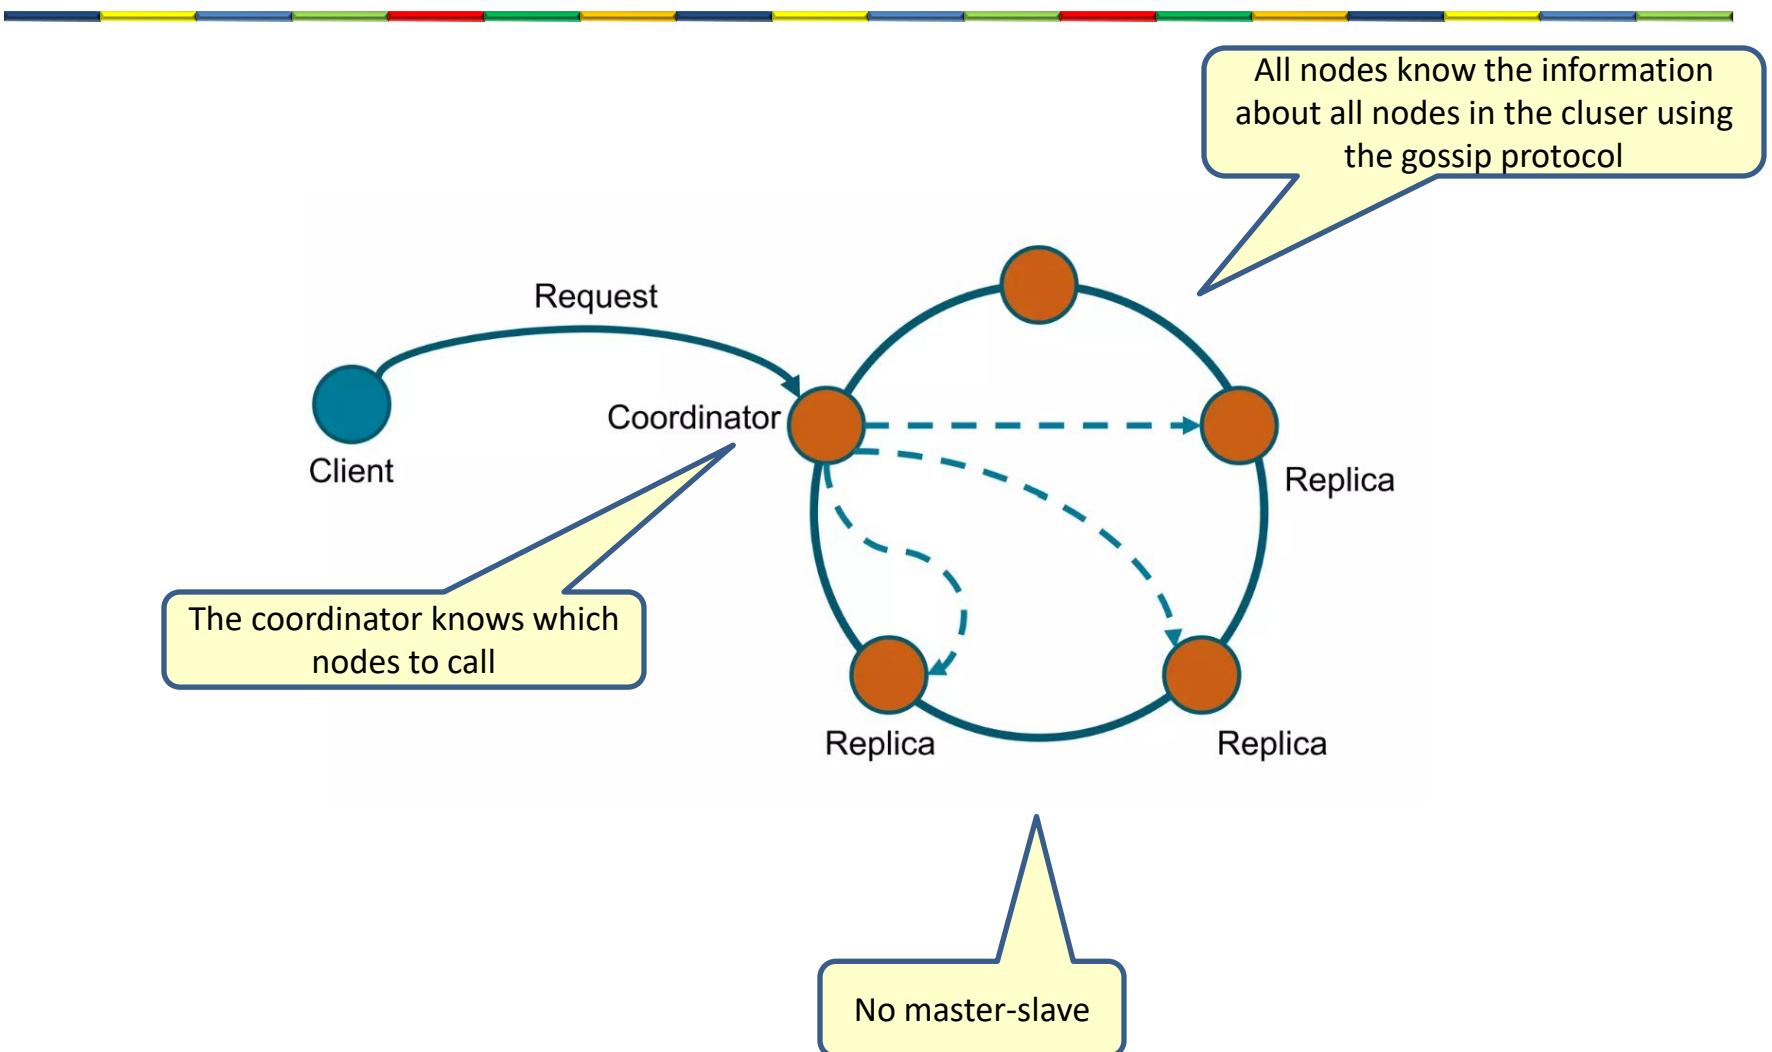
first_name	last_name	title	year
Johnny	Depp	Alice in Wonderland	2010
Johnny	Depp	Edward Scissorhands	1990
Anne	Hathaway	Alice in Wonderland	2010

Cassandra replication

RF = 3

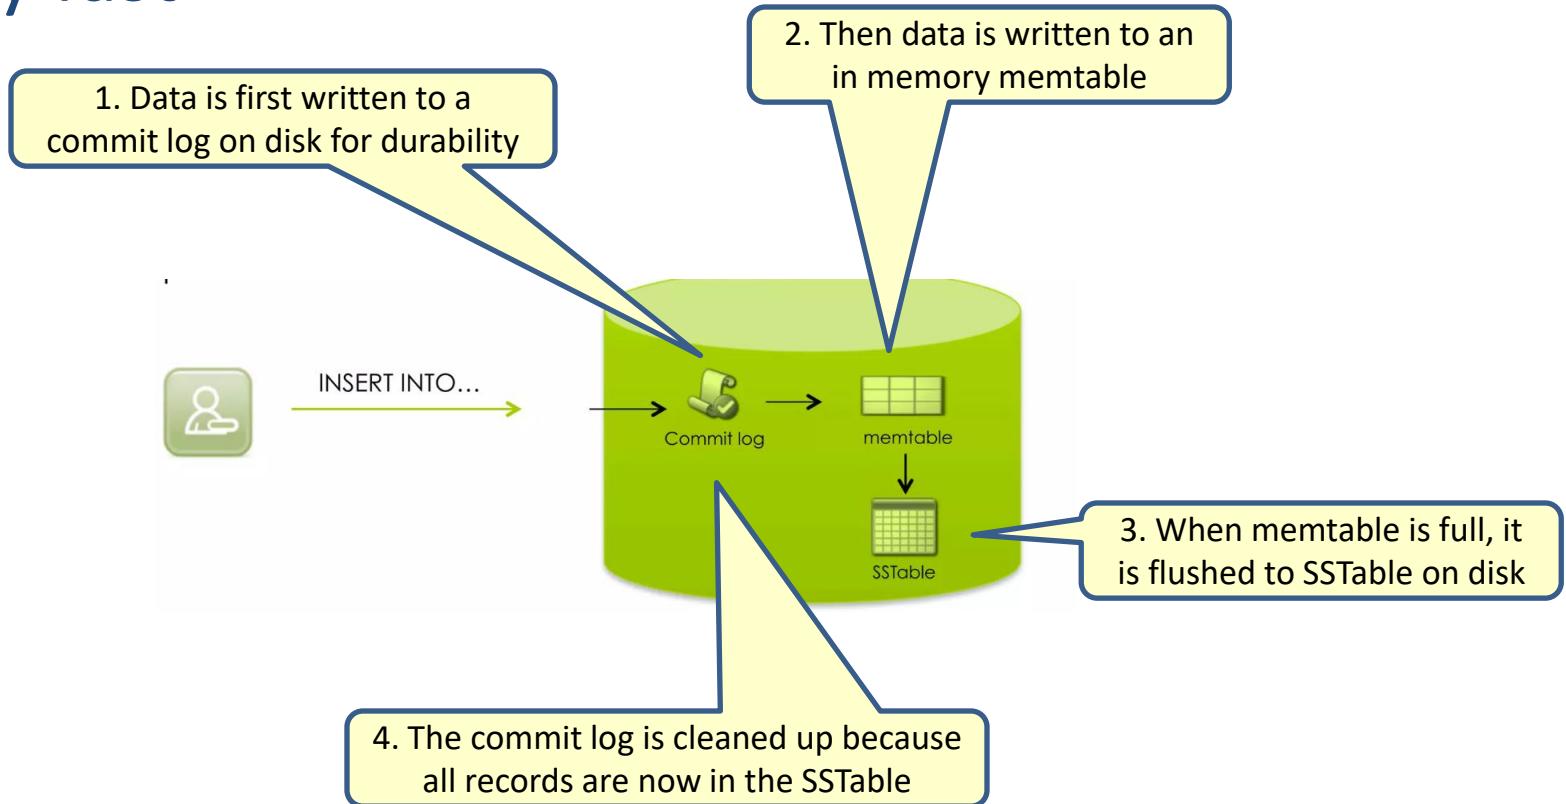


You can read and write to any node



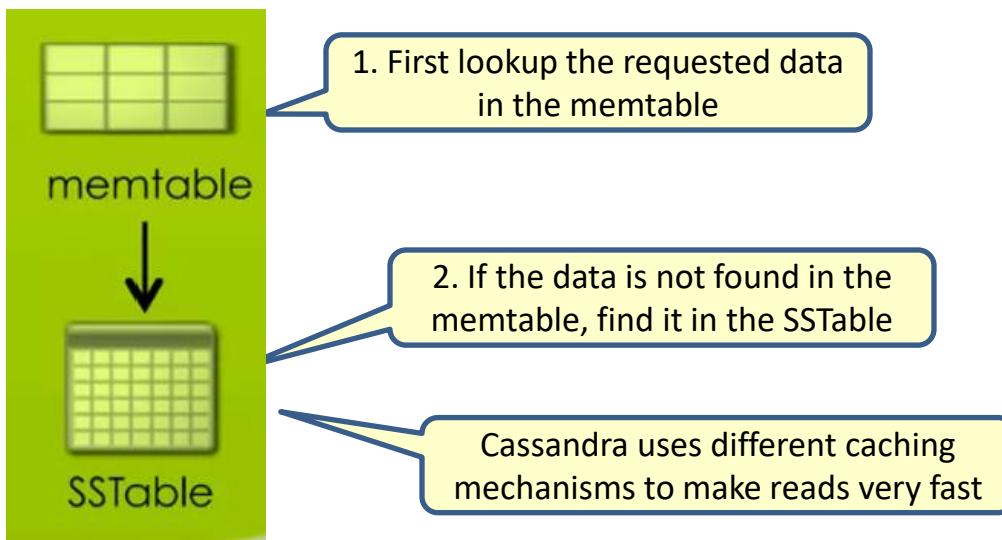
Writes in Cassandra

- Very fast



Reads in Cassandra

- Very fast

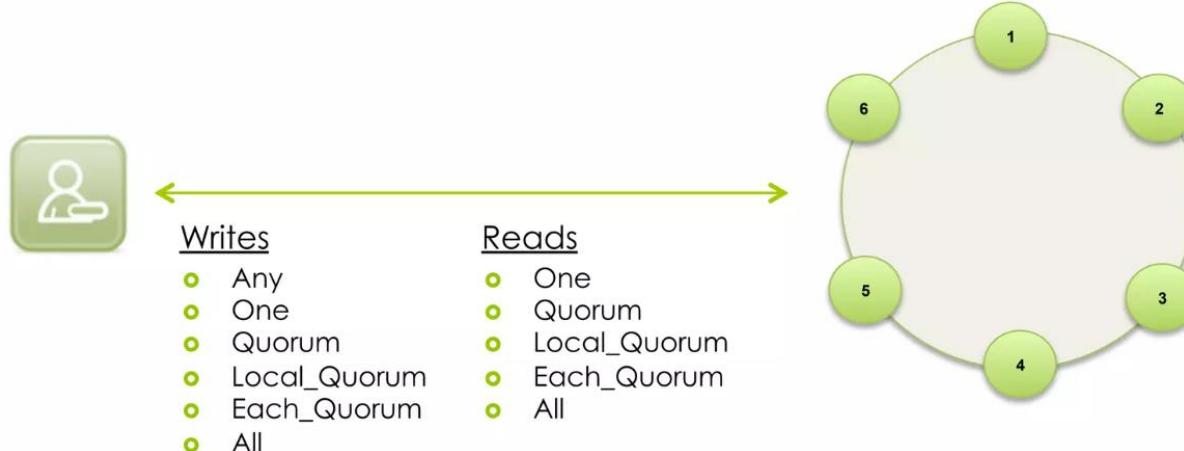


Tunable consistency

- Choose the balance between **data consistency** and **availability** for each read and write operation by setting a consistency level.
- **Write operations:** The client determines how many replicas must acknowledge the write before the operation is reported as successful.
- **Read operations:** The client specifies how many replicas must respond to a read request before data is returned to the client.
- **Key consistency levels**
 - **ONE:** The operation is successful if only one replica responds. This is the fastest but offers the lowest consistency.
 - **QUORUM:** A majority of replicas must respond. For a replication factor of 3, this means 2 out of 3 replicas must acknowledge the operation. This provides a balance between consistency and performance.
 - **ALL:** All replicas must respond. This provides the strongest consistency but is the slowest option.



Cassandra tunable consistency



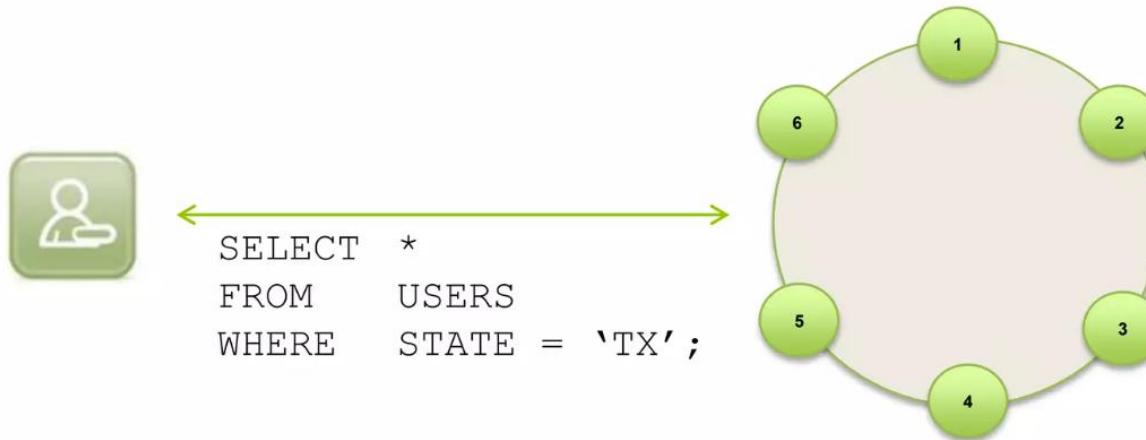
CQL Examples

```
SELECT total_purchases FROM SALES  
USING CONSISTENCY QUORUM  
WHERE customer_id = 5
```

```
UPDATE SALES  
USING CONSISTENCY ONE  
SET total_purchases = 500000  
WHERE customer_id = 4
```

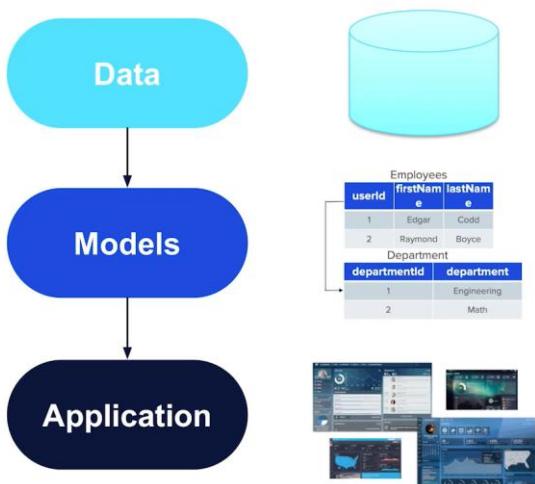
Cassandra Query Language (CQL)

- Very similar to SQL

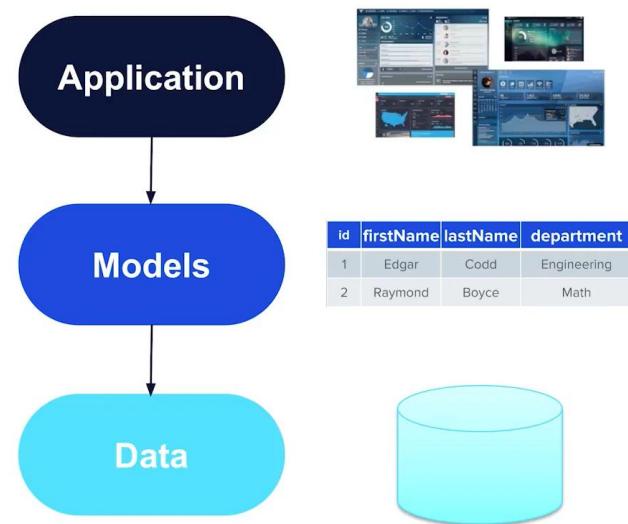


Data modeling

■ RDBMS



• Cassandra



Ratings by movie

title	year	email	rating
Alice in Wonderland	2010	jen@datastax.com	10
Alice in Wonderland	2010	joe@datastax.com	9
Alice in Wonderland	1951	jen@datastax.com	8
Edward Scissorhands	1990	joe@datastax.com	10

```
CREATE TABLE ratings_by_movie (
    title TEXT,
    year INT,
    email TEXT,
    rating INT,
    PRIMARY KEY ((title, year), email)
); ↵
```

```
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 2010,
'jen@datastax.com', 10);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 2010,
'joe@datastax.com', 9);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Alice in Wonderland', 1951,
'jen@datastax.com', 8);
INSERT INTO ratings_by_movie (title, year, email,
rating)
VALUES ('Edward Scissorhands', 1990,
'joe@datastax.com', 10); ↵
```

Actors by movie



title	year	first_name	last_name
Alice in Wonderland	2010	Anne	Hathaway
Alice in Wonderland	2010	Johnny	Depp
Edward Scissorhands	1990	Johnny	Depp

```
CREATE TABLE actors_by_movie (
    title TEXT,
    year INT,
    first_name TEXT,
    last_name TEXT,
    PRIMARY KEY ((title, year), first_name, last_name)
); ↵
```



Movies by actor

first_name	last_name	title	year
Johnny	Depp	Alice in Wonderland	2010
Johnny	Depp	Edward Scissorhands	1990
Anne	Hathaway	Alice in Wonderland	2010

```
CREATE TABLE movies_by_actor (
    first_name TEXT,
    last_name TEXT,
    title TEXT,
    year INT,
    PRIMARY KEY ((first_name, last_name), title, year)
); ↵
```

Movies by user

email	watched_on	title	year
joe@datastax.com	2020-04-28	Edward Scissorhands	1990
joe@datastax.com	2020-03-08	Alice in Wonderland	2010
joe@datastax.com	2020-02-13	Toy Story 3	2010
joe@datastax.com	2020-01-22	Despicable Me	2010
joe@datastax.com	2019-12-30	Alice in Wonderland	1951
jen@datastax.com	2011-10-01	Alice in Wonderland	2010

```
CREATE TABLE movies_by_user (
    email TEXT,
    title TEXT,
    year INT,
    watched_on DATE,
    PRIMARY KEY ((email), watched_on, title, year)
) WITH CLUSTERING ORDER BY (watched_on DESC); ↵
```



Data duplication



email	title	year	rating
joe@datastax.com	Alice in Wonderland	2010	9
joe@datastax.com	Edward Scissorhands	1990	10
jen@datastax.com	Alice in Wonderland	1951	8
jen@datastax.com	Alice in Wonderland	2010	10

title	year	first_name	last_name
Alice in Wonderland	2010	Anne	Hathaway
Alice in Wonderland	2010	Johnny	Depp
Edward Scissorhands	1990	Johnny	Depp

title	year	email	rating
Alice in Wonderland	2010	jen@datastax.com	10
Alice in Wonderland	2010	joe@datastax.com	9
Alice in Wonderland	1951	jen@datastax.com	8
Edward Scissorhands	1990	joe@datastax.com	10

email	watched_on	title	year
joe@datastax.com	2020-04-28	Edward Scissorhands	1990
joe@datastax.com	2020-03-08	Alice in Wonderland	2010
joe@datastax.com	2020-02-13	Toy Story 3	2010
joe@datastax.com	2020-01-22	Despicable Me	2010
joe@datastax.com	2019-12-30	Alice in Wonderland	1951
jen@datastax.com	2011-10-01	Alice in Wonderland	2010



Cassandra advantages & disadvantages

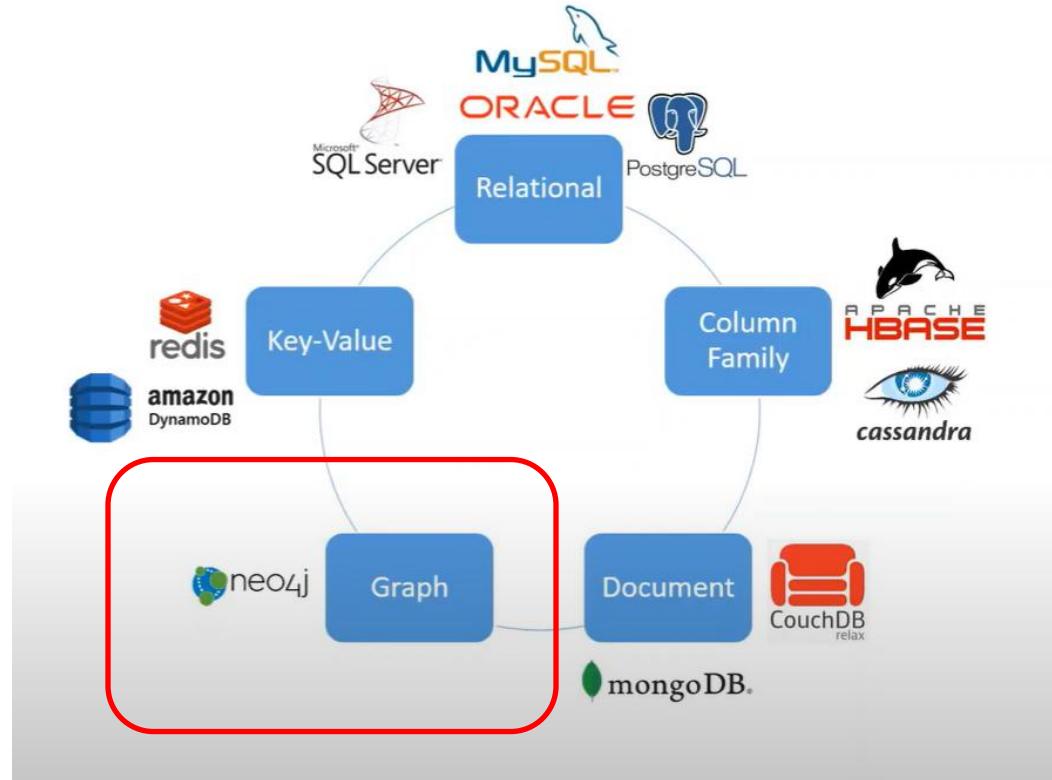
- Advantages
 - Very fast for writes and reads
 - Linear scalability
 - You can write to any node to geographically distributed clusters
 - Tunable consistency
- Disadvantages
 - No joins and aggregates
 - No transactions
 - No strict consistency
 - Data duplication
 - Slow updates and deletes
 - Only structured data



Cassandra use cases

- Real-time analytics pipeline
 - Large data sets, only write and read
- IoT data management
 - Large data sets, only write and read
- Content management systems
 - High write volumes, ensure data availability, and maintain consistency
- Real-time analysis of large datasets
 - Fraud detection and risk management
- Time-series data analysis
 - Large data sets, only write and read





GRAPH DATABASE



NEO4J



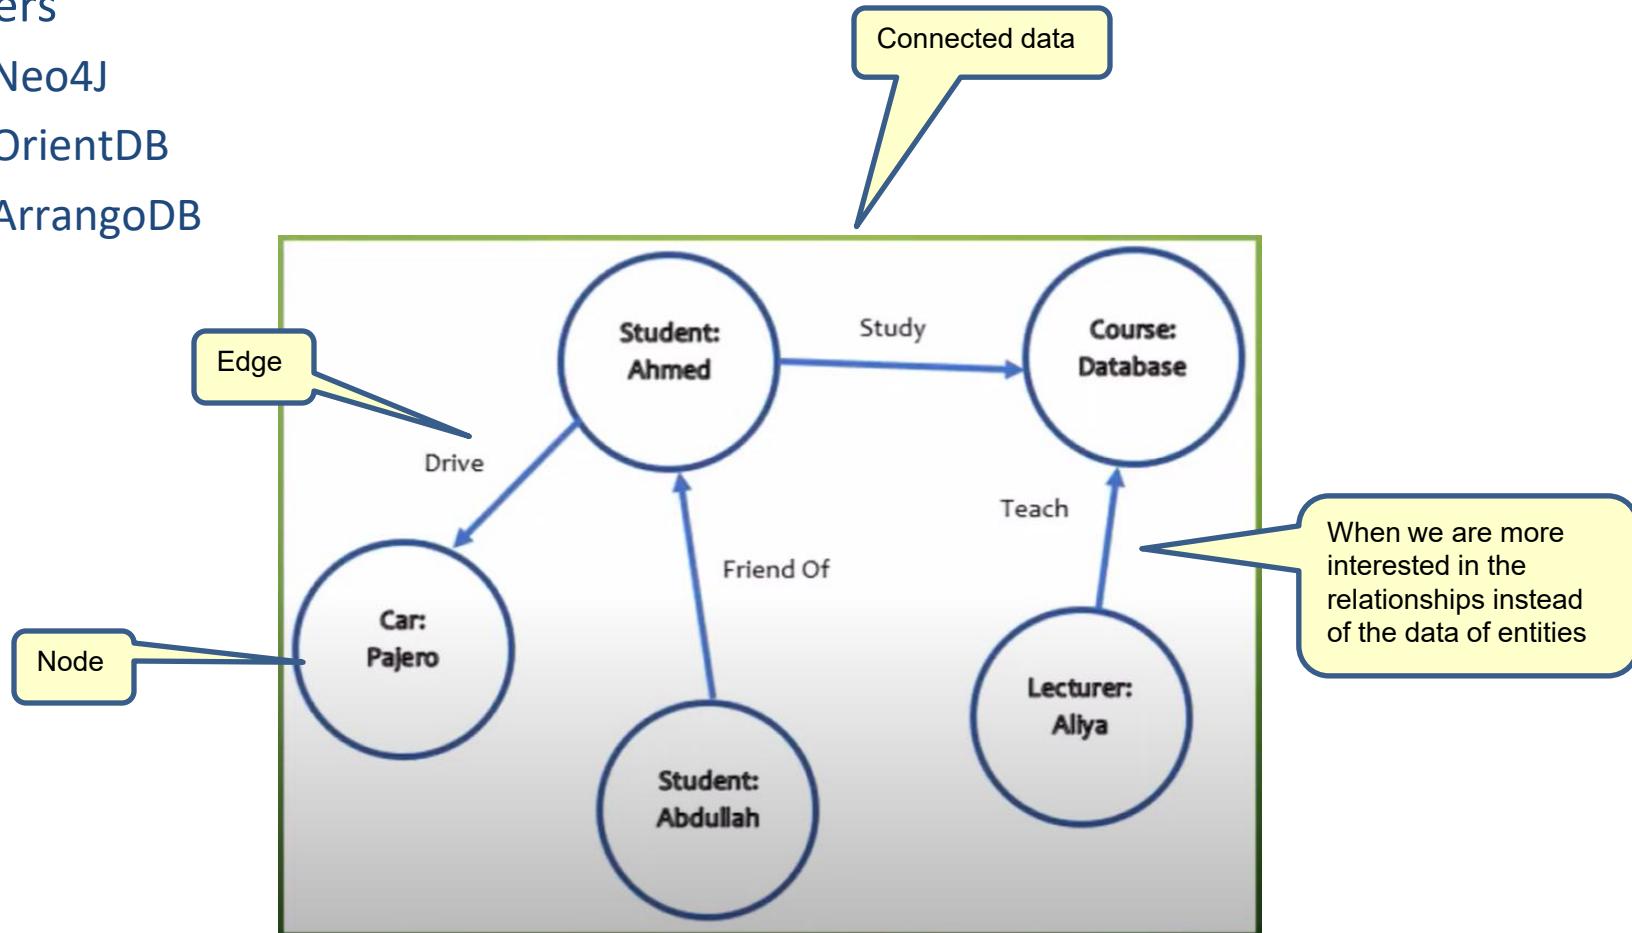
Graph database

- Store data in nodes and relationships
 - No Joins
- Very fast in analyzing data with lots of relationships
- No fixed schema
 - Easy to evolve your dataset



Graph database

- Players
 - Neo4J
 - OrientDB
 - ArangoDB

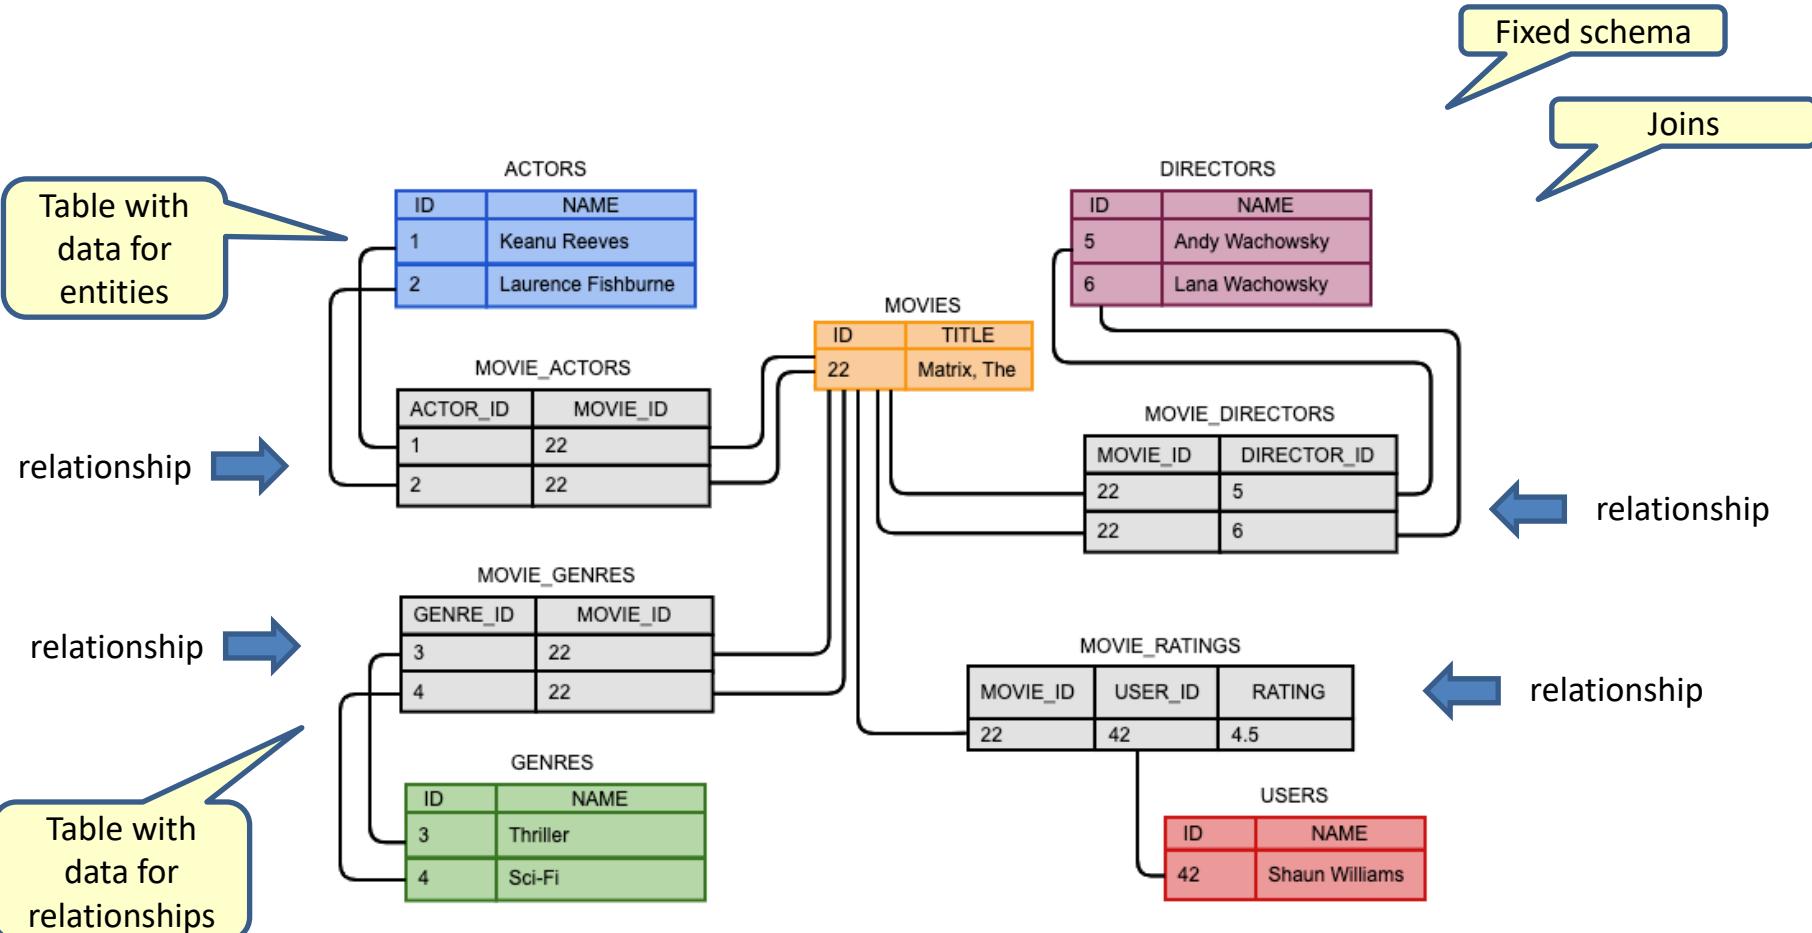


Graph DB use cases

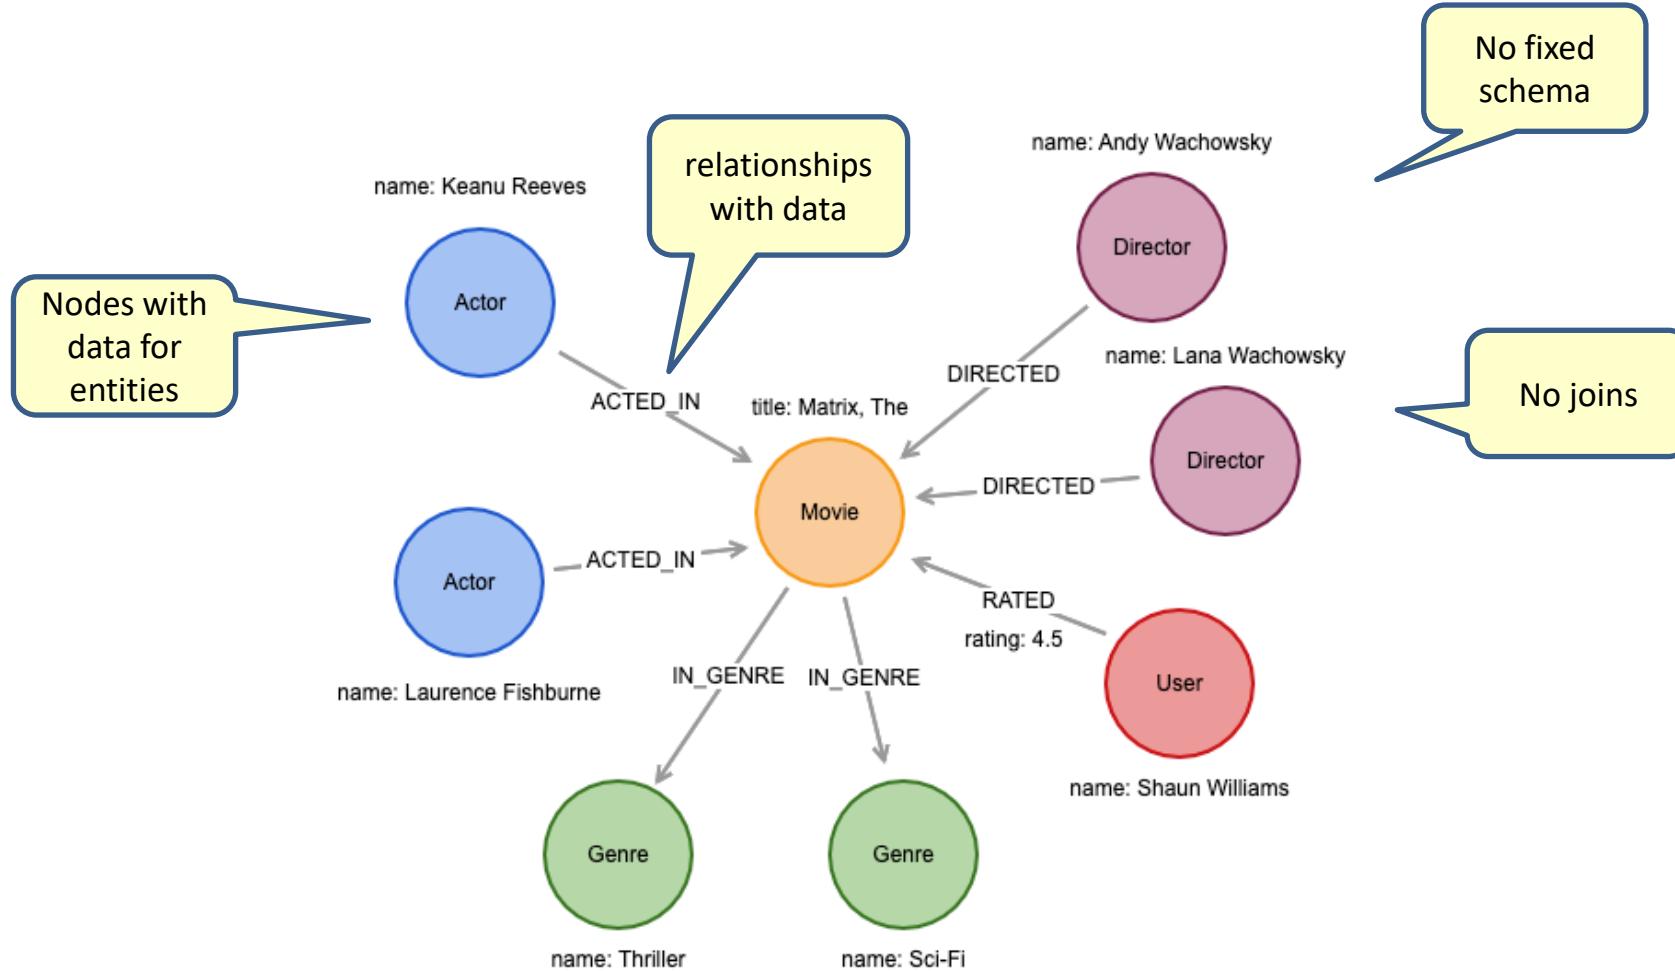
- Social networks
- Knowledge graphs
- Fraud detection
- Contact tracing
- Real time recommendation
- Graph based search
- Network & IT operations
- Identity and access management



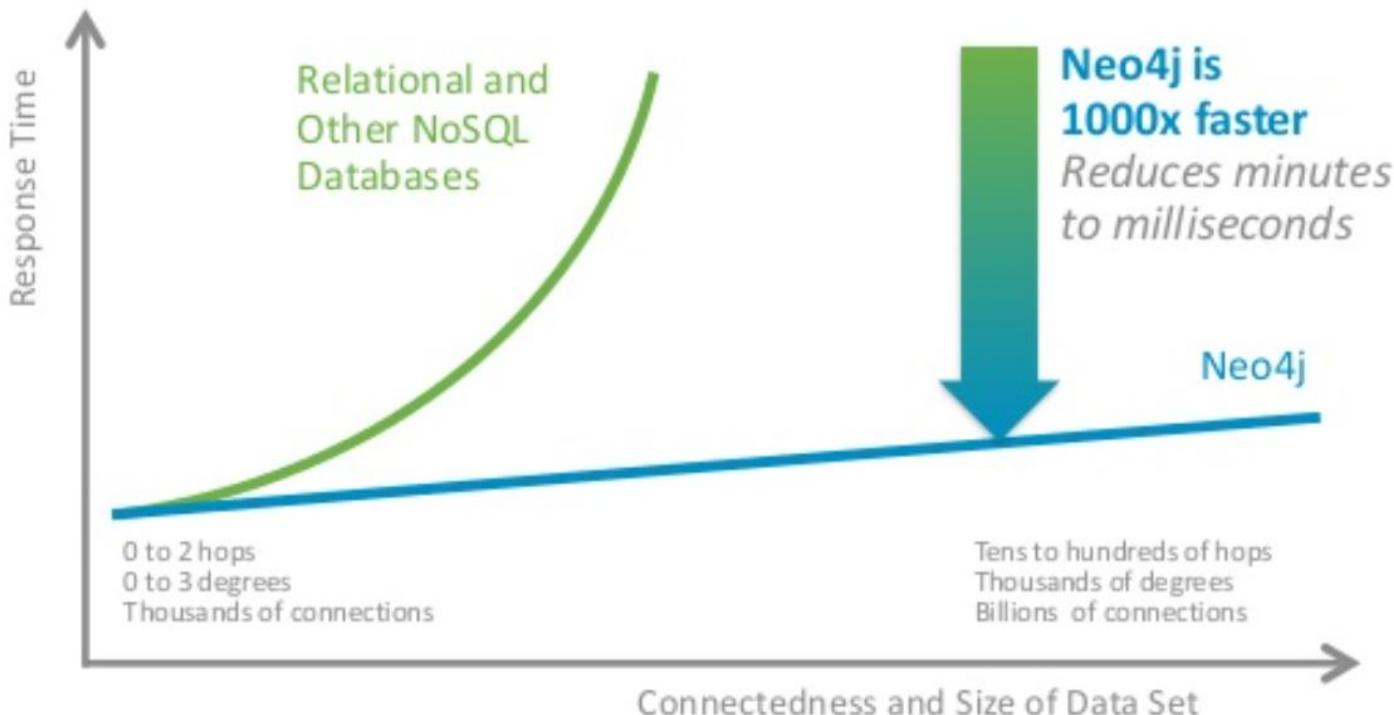
Relational database



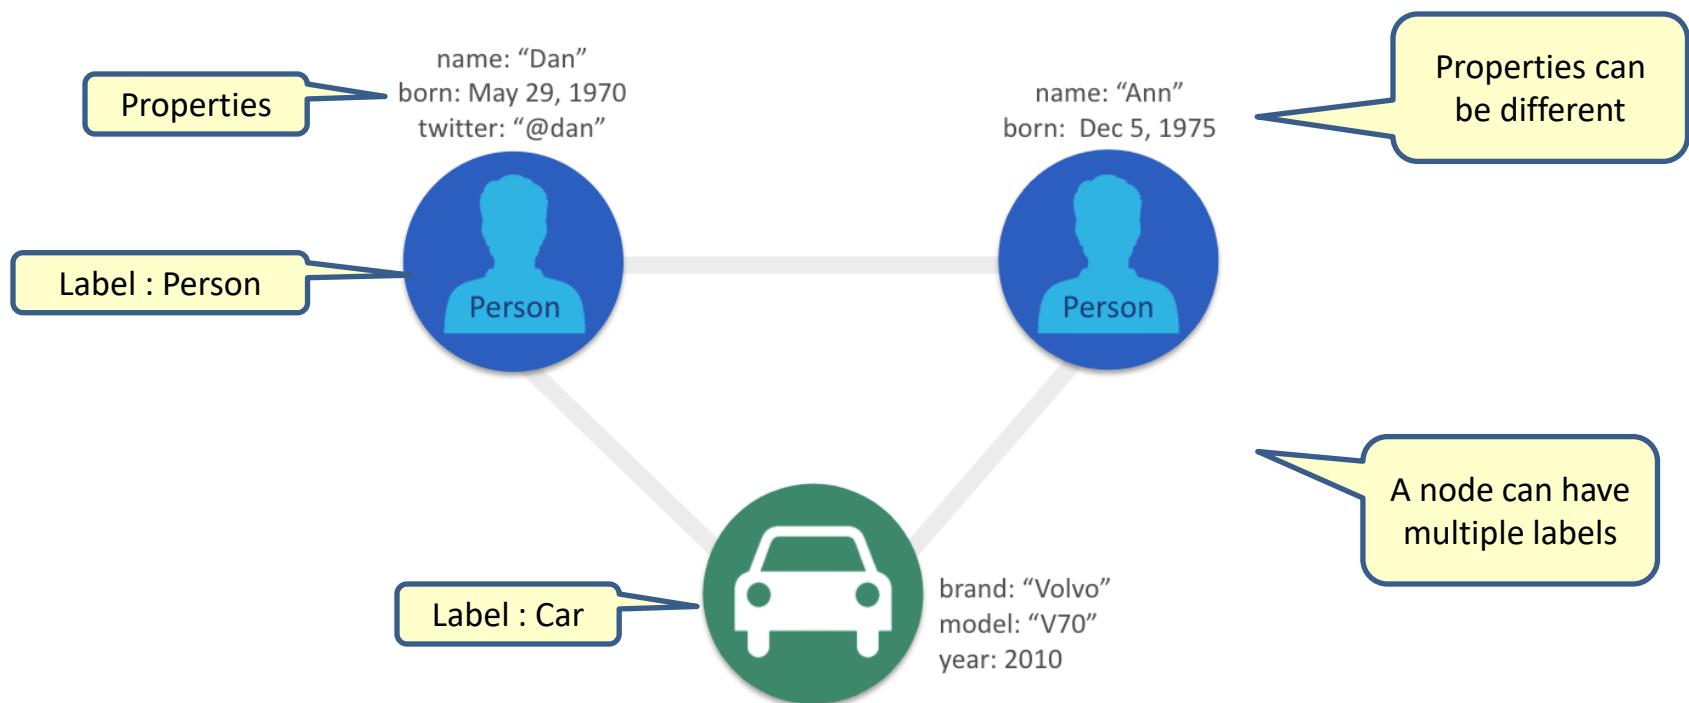
Graph database: NEO4J



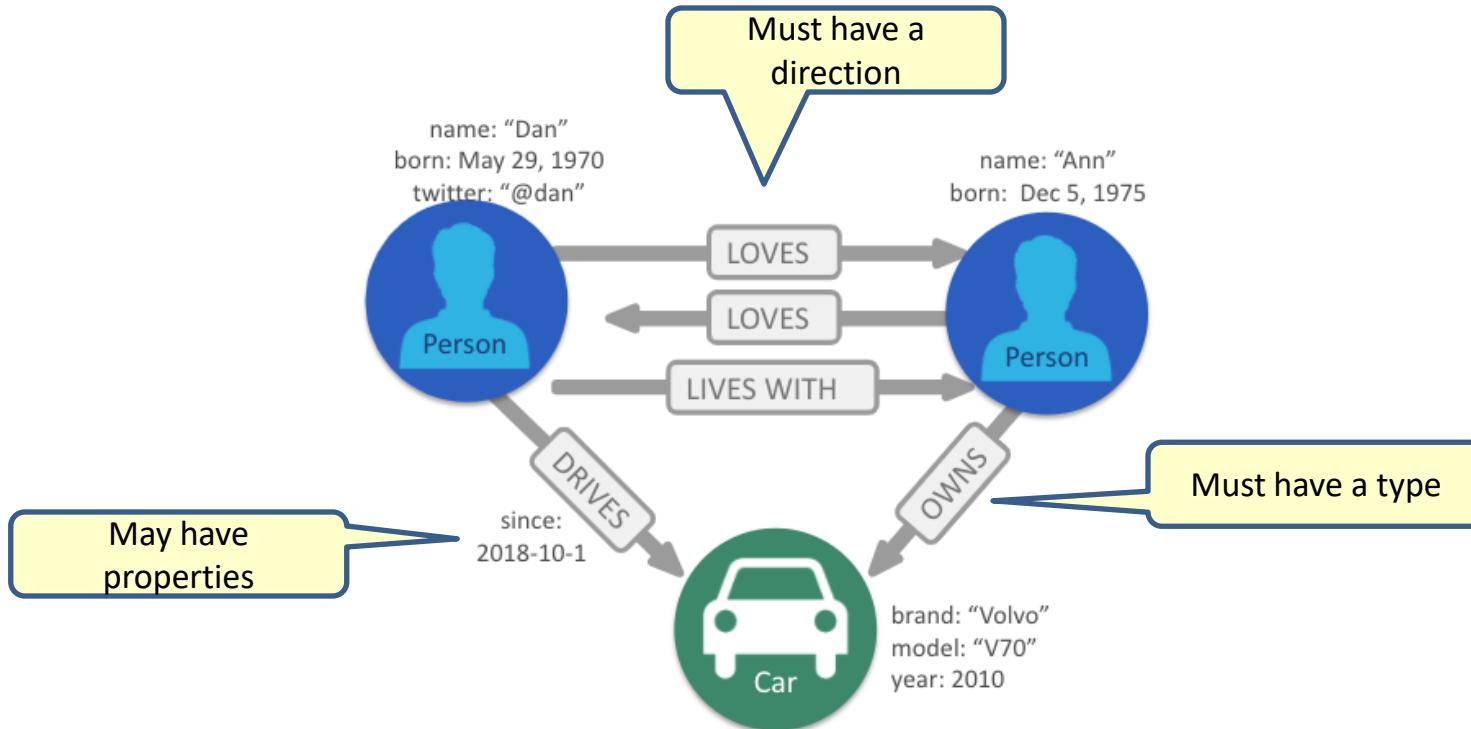
Query performance



Nodes

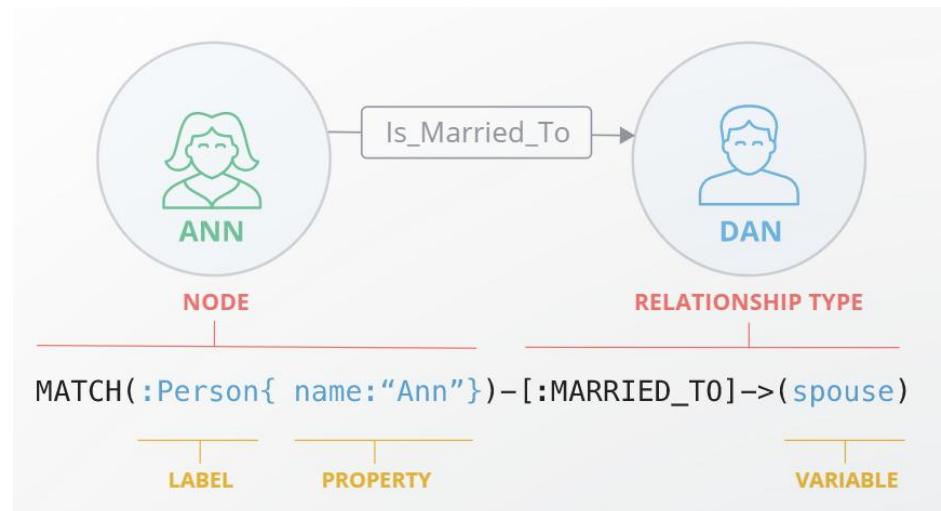


Relationships

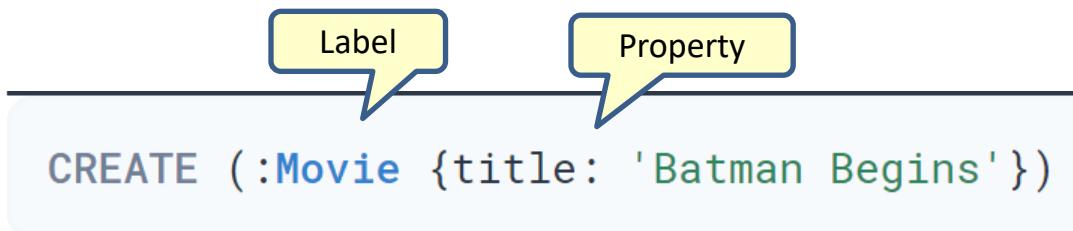


Cypher query language

- Declarative
 - Describe what you want
 - Not how to do it



Create a node



The screenshot shows the Neo4j browser interface. On the left, a sidebar has icons for Graph, Table (selected), Text, and Code. The main area displays the results of a query:

```
neo4j$ MATCH (n) RETURN n
```

Table output:

n
1
{ "identity": 0, "labels": ["Movie"], "properties": { "title": "Batman Begins" } }

Graph output:

```
neo4j$ MATCH (n) RETURN n
```

Shows a single node represented by an orange circle with the text "Batman Begins" inside. To the left of the node, there are two status indicators: a purple oval labeled "*1" and an orange oval labeled "Movie(1)".

Create a relationship

```
MATCH (a:Person), (m:Movie)
WHERE a.name = 'Michael Caine' AND m.title = 'Batman Begins'
CREATE (a)-[:ACTED_IN]->(m)
```



```
MATCH (a:Person), (m:Movie)
WHERE a.name = 'Christian Bale' AND m.title = 'Batman Begins'
CREATE (a)-[:ACTED_IN {roles: ['Bruce Wayne', 'Batman']}]->(m)
RETURN a, m
```



Queries on nodes

```
MATCH (n)  
RETURN n
```

Retrieve all nodes

```
MATCH (p:Person)  
RETURN p
```

Retrieve all Person nodes

```
MATCH (p:Person {born: 1970})  
RETURN p
```

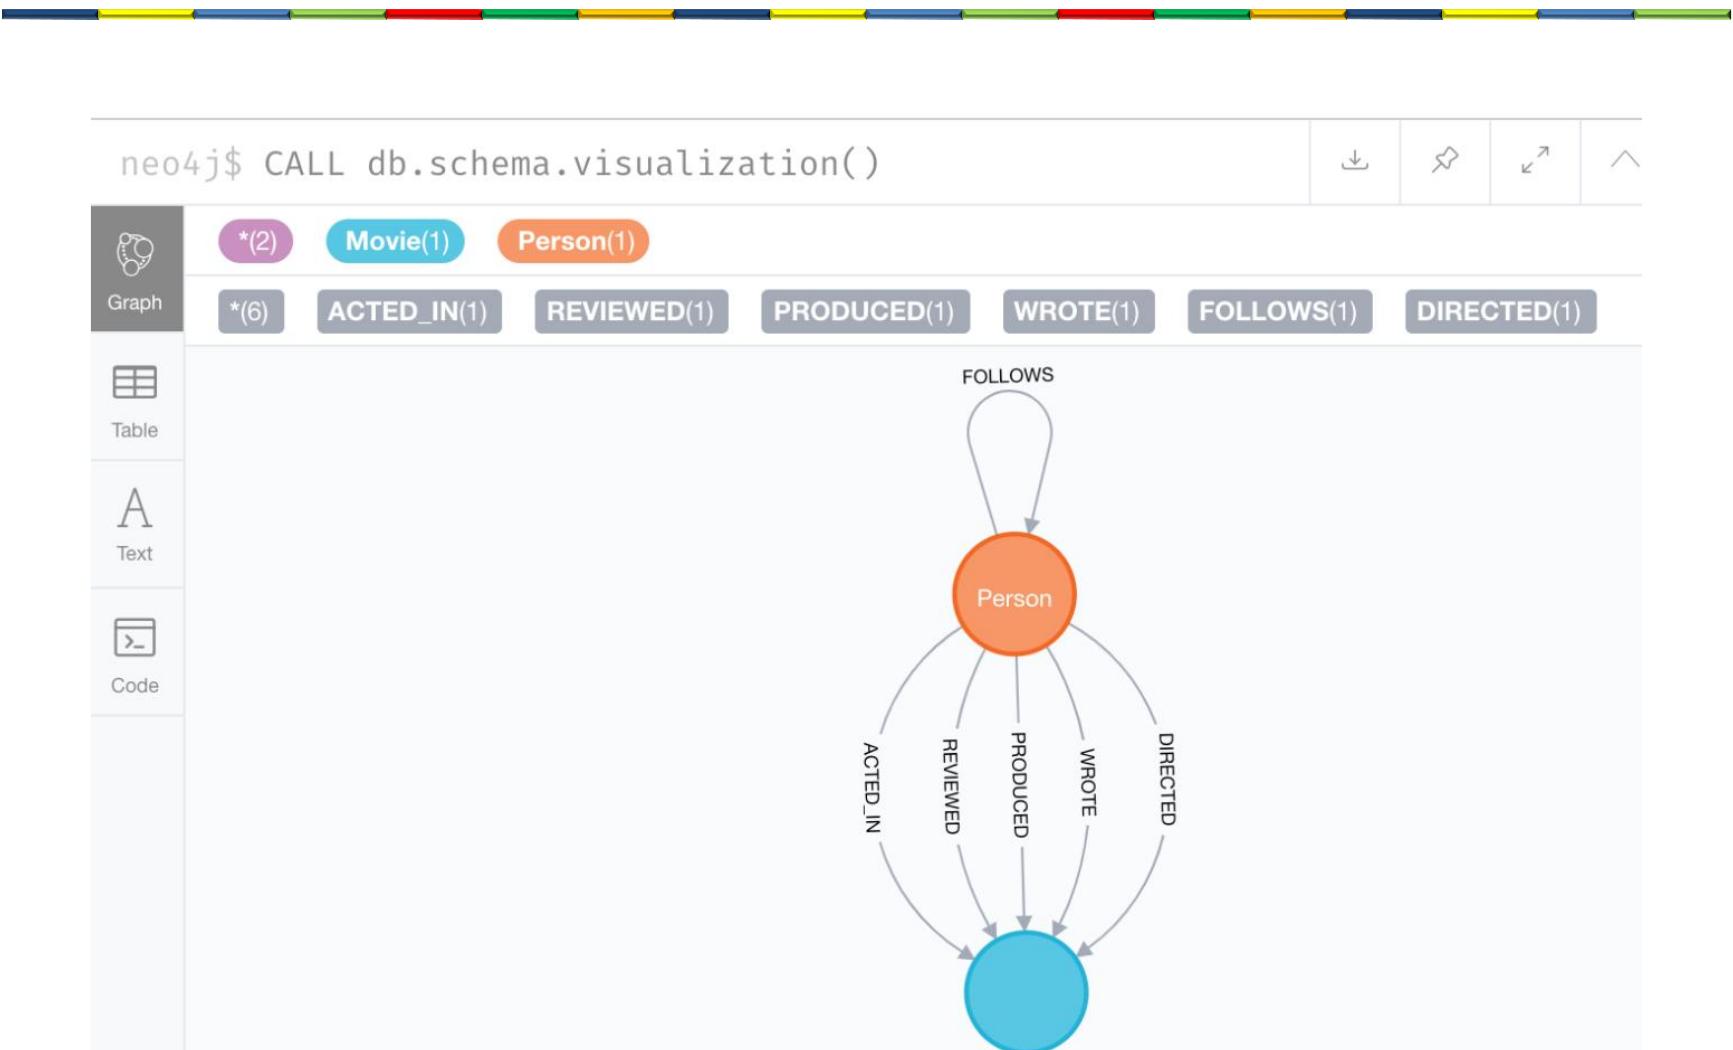
Retrieve all Person nodes born in 1970

```
MATCH (m:Movie {released: 2003, tagline: 'Free your mind'})  
RETURN m
```

Retrieve all Movie nodes released in 2003
with the tagline ‘free your mind’

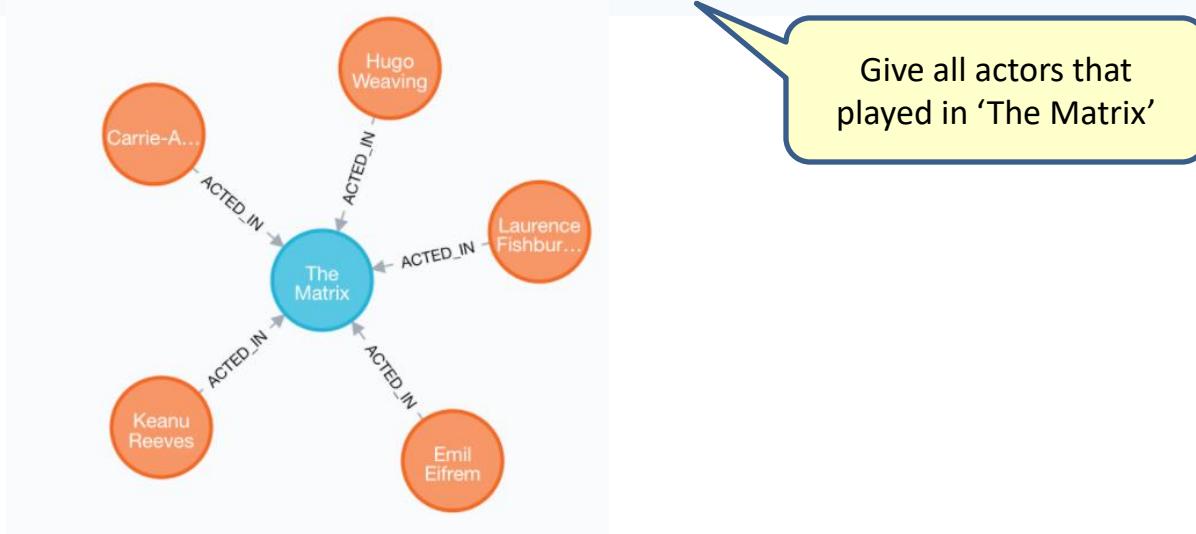


Visualize the schema



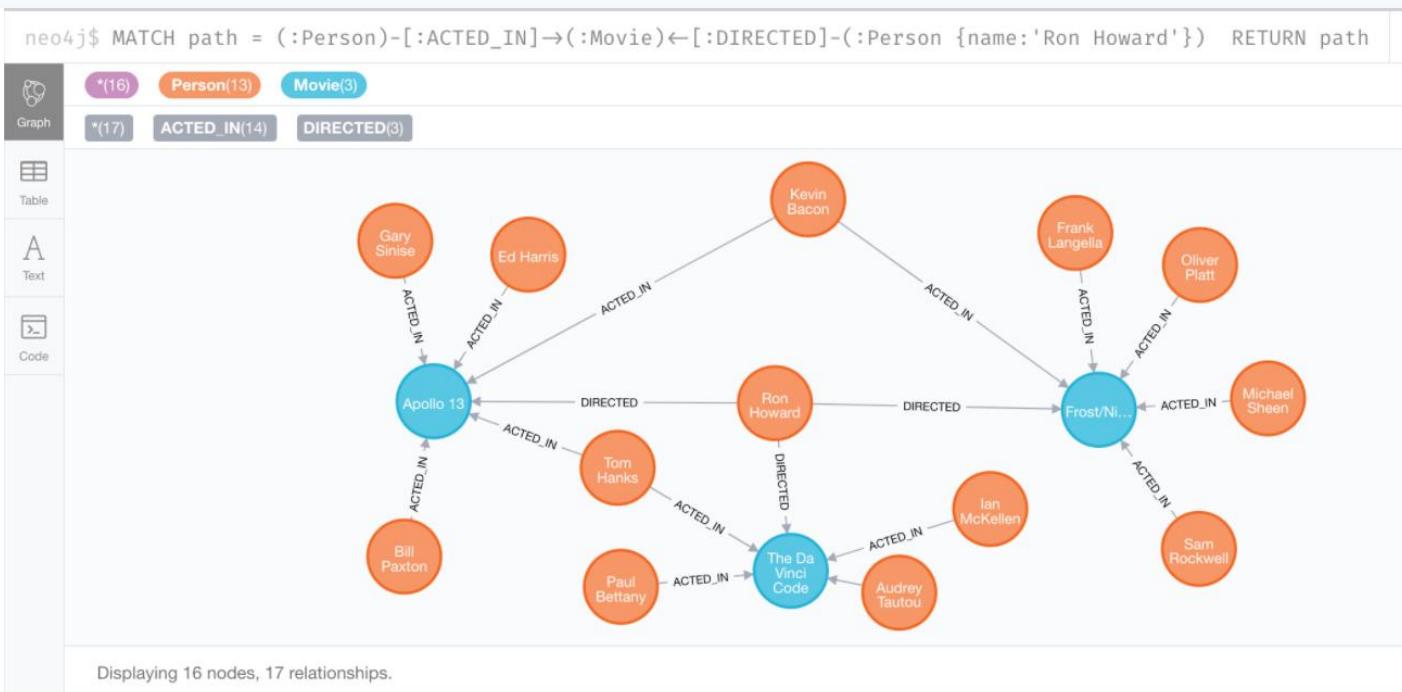
Queries on relationships

```
MATCH (p:Person)-[rel:ACTED_IN]->(m:Movie {title: 'The Matrix'})  
RETURN p, rel, m
```

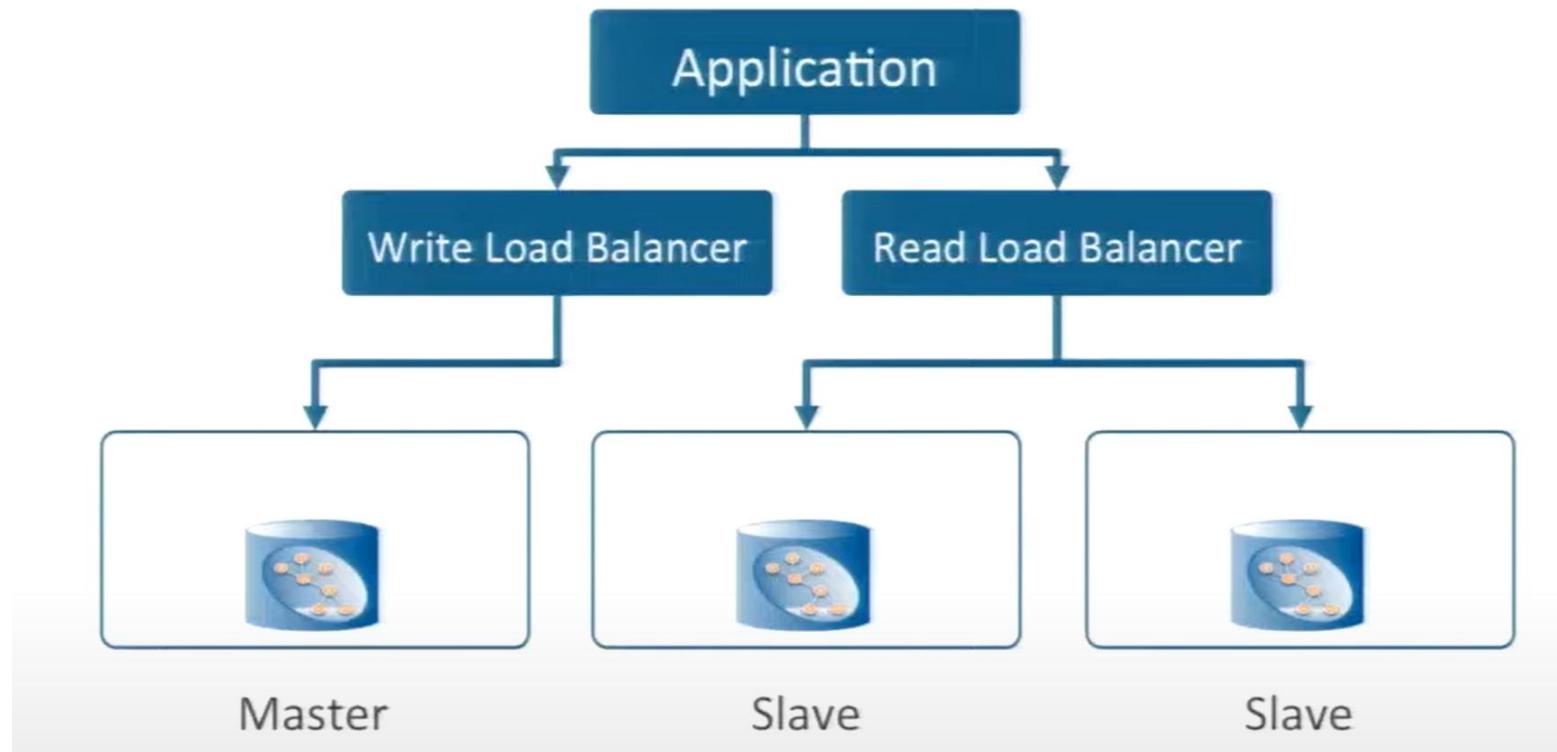


Return multiple paths

```
MATCH path = (:Person)-[:ACTED_IN]->(:Movie)<-[ :DIRECTED] -(:Person {name:'Ron Howard'})  
RETURN path
```



Neo4j High Availability cluster



Neo4j advantages & disadvantages

- Advantages
 - Very fast in querying deeply connected data
 - Database model and domain model are similar
 - Powerful and simple query language
 - Easy to add new nodes, attributes and relationships
 - Cluster (scalability, failover and high availability)
- Disadvantages
 - Not for unconnected data
 - Not for unstructured data



SPRING BOOT NEO4J EXAMPLE



Domain classes + repository

```
@Node  
public class Person {  
  
    @Id @GeneratedValue private Long id;  
  
    private String name;  
  
    @Relationship(type = "TEAMMATE")  
    public Set<Person> teammates= new HashSet<>();
```

```
public interface PersonRepository extends Neo4jRepository<Person, Long> {  
  
    Person findByName(String name);  
    List<Person> findByTeammatesName(String name);  
}
```



Application(1/2)

```
@SpringBootApplication
@EnableNeo4jRepositories
public class PersonApplication implements CommandLineRunner{

    @Autowired
    PersonRepository personRepository;

    public static void main(String[] args) {
        SpringApplication.run(PersonApplication.class, args);
    }
}
```



Application(2/2)

```
@Override
public void run(String... args) throws Exception {
    personRepository.deleteAll();

    Person greg = new Person("Greg");
    Person roy = new Person("Roy");
    Person bob = new Person("Bob");
    greg.worksWith(roy);
    greg.worksWith(bob);
    roy.worksWith(bob);

    List<Person> team = Arrays.asList(greg, roy, bob);

    personRepository.save(greg);
    personRepository.save(roy);
    personRepository.save(bob);

    System.out.println("Lookup each person by name...");
    team.stream().forEach(person -> System.out.println(
        personRepository.findByName(person.getName())));
}

List<Person> teammates = personRepository.findByTeammatesName(bob.getName());
System.out.println("The following have Bob as a teammate...");
teammates.stream().forEach(person -> System.out.println( person.getName()));

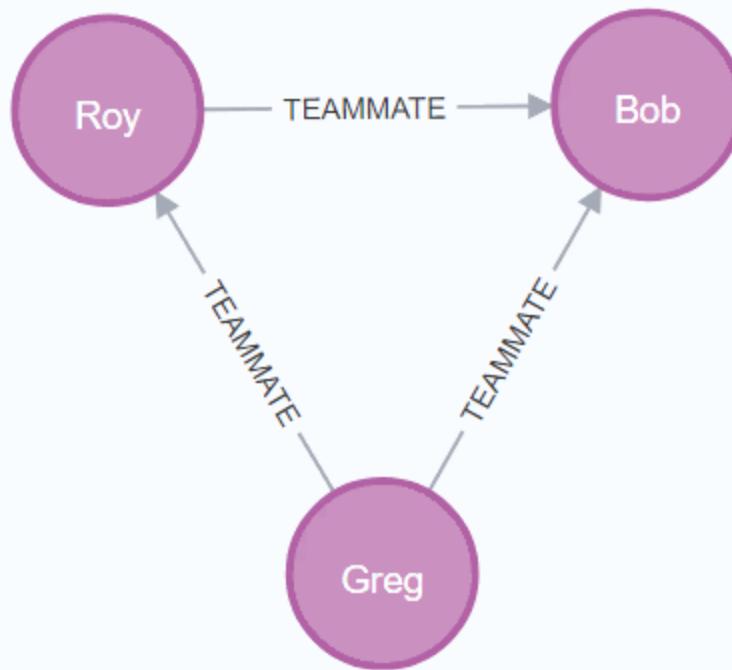
}
```

application.properties

```
spring.neo4j.uri=bolt://localhost:7687  
spring.data.neo4j.username=neo4j  
spring.data.neo4j.password=admin
```



Nodes



SUMMARY

Five different types of databases

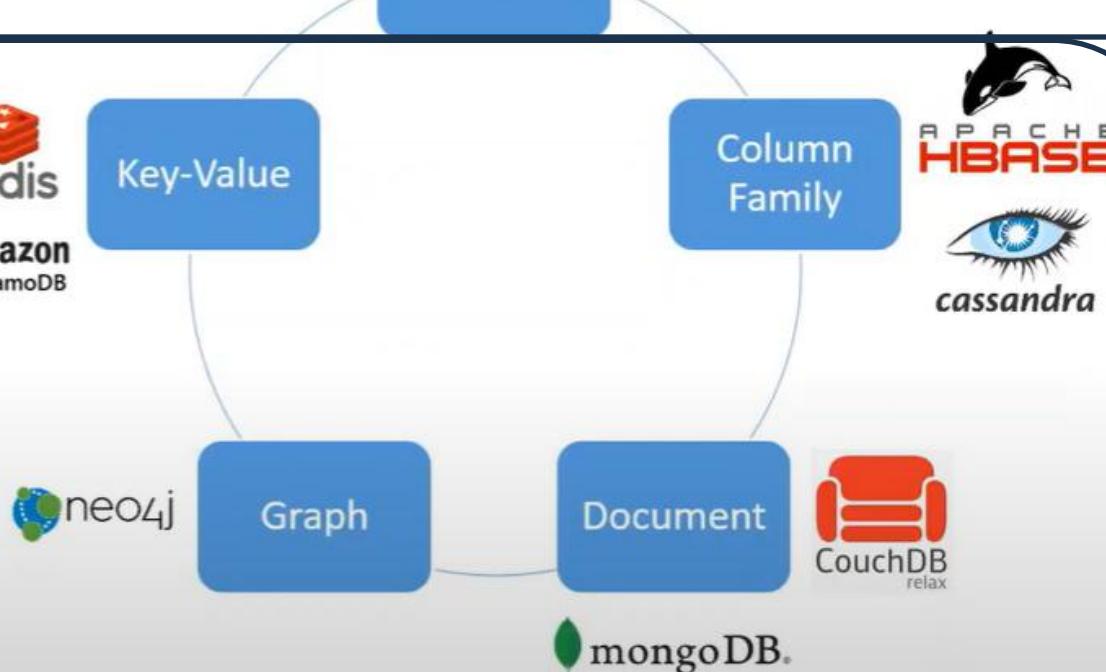
Relational



Non-relational



NoSQL
(Not only SQL)



Relational database

- Does not scale well
- Hard to change
- Does not handle unstructured and semi-structured data



NoSQL databases

- Scales very well
- Easy to change, no fixed schema
- Handles unstructured and semi structured data



Comparing databases

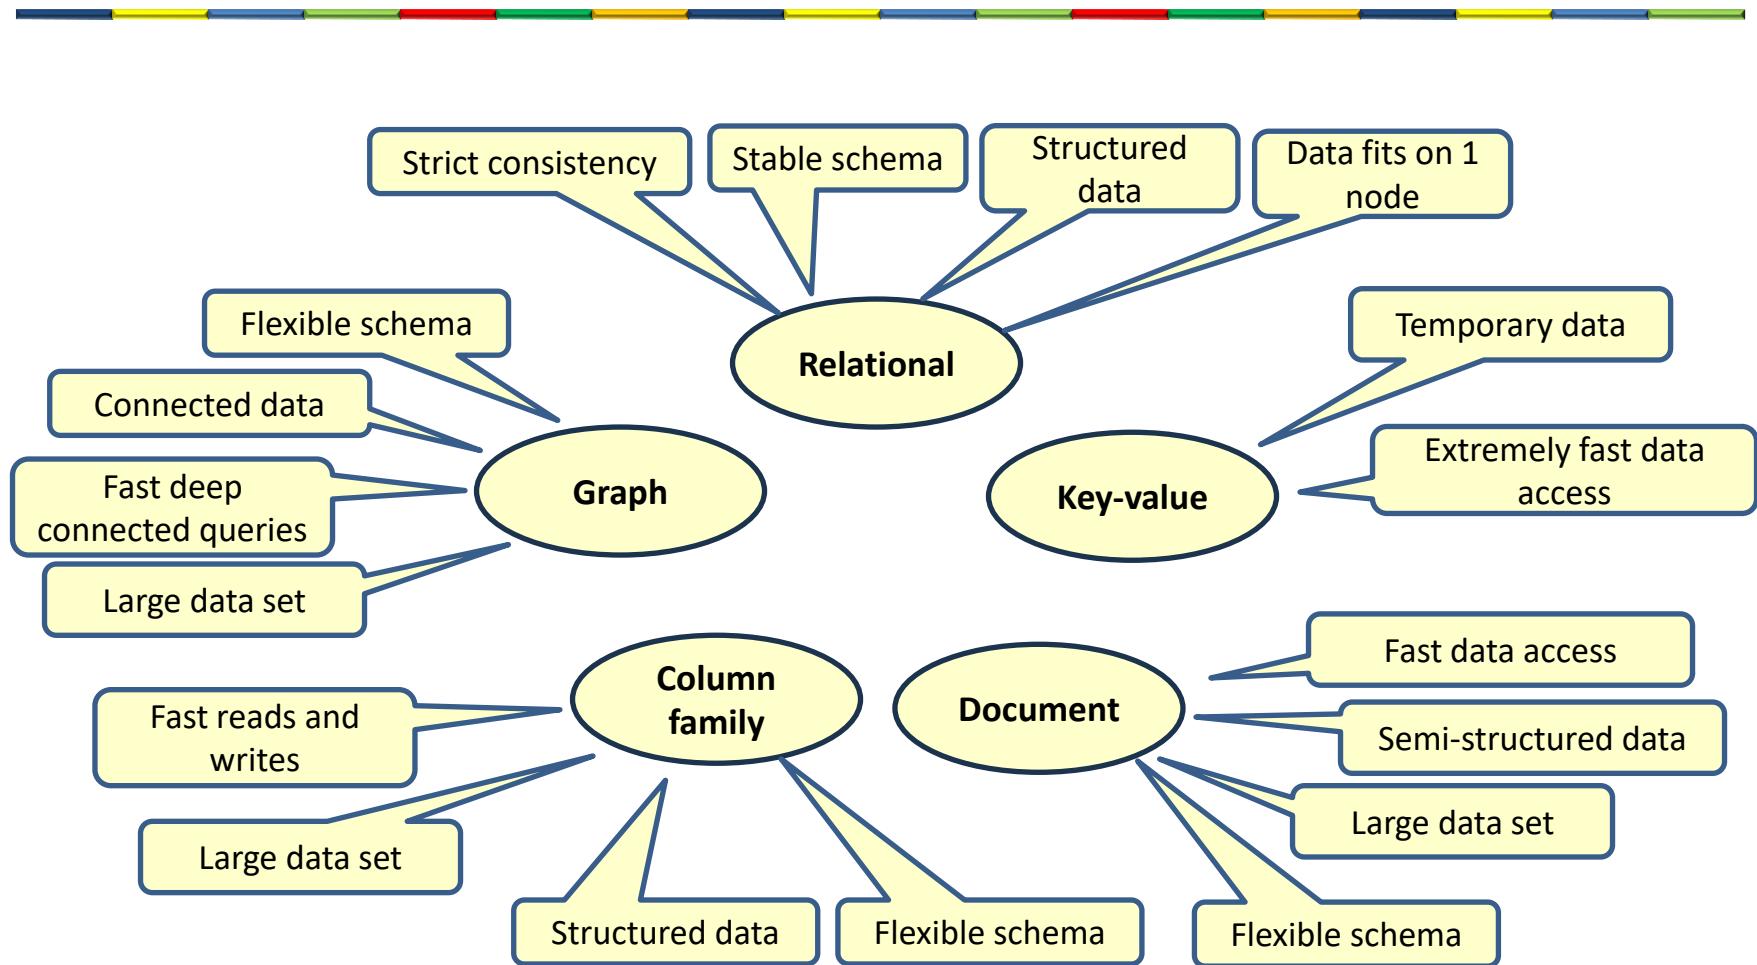
		Strong points	Weak points
Relational	Structured data Wide adaption Strict consistency	Hard to scale Hard to change the schema Not good in unstructured data	
Key-value	Extremely fast in-memory data	Not for permanent data storage No or limited query functionality	
Document	Whole documents Fast reads and writes Unstructured data	Data duplication	
Wide column	Very fast reads and writes Fast queries at scale Linear scalable	Data duplication Slow updates and deletes Not good in unstructured data	
Graph	Connected data Fast queries at scale Flexible data structure	Not good for unconnected data Not good in unstructured data	

Transactions

- Oracle:
 - One node: ACID
 - HA cluster with sharding: ACID
- Mongo, Neo4J:
 - One node: ACID
 - HA cluster with sharding: BASE
- Cassandra:
 - One node: Tunable consistency
 - HA cluster with sharding: Tunable consistency



Which database to use?



Key principle 6

- When your data does not fit on one node you automatically will get
 - Data duplication
- Data duplication over multiple nodes means eventual consistency

